


Introduction to Spark

Christoph Kessler
IDA, Linköping University

Christoph Kessler, IDA,
Linköpings universitet.

Recall: MapReduce Programming Model

- Designed to operate on LARGE distributed input data sets stored e.g. in HDFS nodes
- Abstracts from parallelism, data distribution, load balancing, data transfer, fault tolerance
- Implemented in **Hadoop** and other frameworks
- Provides a high-level parallel programming construct (= a skeleton) called **MapReduce**
 - A generalization of the data-parallel *MapReduce* skeleton of Lect. 1
 - Covers the following algorithmic design pattern:



C. Kessler, IDA, Linköpings universitet.

2

From MapReduce to Spark

MapReduce

- is for large-scale computations matching the *MapReduce* pattern,
- with input, intermediate and output data stored in secondary storage

Limitations

- For complex computations composed of *multiple* MapReduce steps
 - E.g. iterative computations
 - e.g. parameter optimization by gradient search
- Much unnecessary disk I/O – data for next MapReduce step could remain in main memory or even cache memory
- Data blocks used multiple times are read multiple times from disk
- Bad data locality across subsequent Mapreduce phases
- Sharing of data only in secondary storage
 - Latency can be too long for interactive analytics
- Fault tolerance by replication of data – more I/O to store copies → slow

By chaining multiple MapReduce steps, we can emulate any distributed computation.

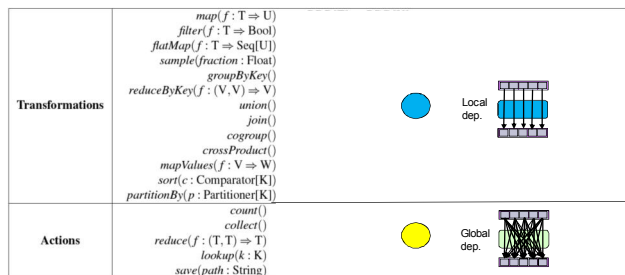


C. Kessler, IDA, Linköpings universitet.

3

Splitting the MapReduce Construct into Simpler Operations – 2 Main Categories:

- Transformations:** Elementwise operations, fully parallelizable
 - Mostly variants of **Map** and reading from distributed file
- Actions:** Operations with internally global dependence structure
 - Mostly variants of **Reduce** and writing back to non-distr. file / to master



RDD transformations and actions available in Spark. Seq[T] denotes a sequence of elements of type T.
Table source: Zaharia et al.: Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing



C. Kessler, IDA, Linköpings universitet.

4

Spark Idea: Data Flow Computing in Memory

Instead of calling subsequent rigid MapReduce steps, the Spark programmer describes the overall **data flow graph** of how to compute all intermediate and final results from the initial input data

- And then "pushes the button" for computing (= *materializing* the results) according to the data flow graph
→ *Lazy evaluation*
- More like declarative, functional programming
- Gives more flexibility to the scheduler
 - Better data locality
 - Keep data in memory as capacity permits, can skip unnecessary disk storage of temporary data
- No replication of data blocks for fault tolerance - in case of task failure (worker failure), **recompute** it from available, earlier computed data blocks according to the data flow graph
 - Needs a **container data structure** for operand data that "knows" how its data blocks are to be computed: the **RDD**




C. Kessler, IDA, Linköpings universitet.

5

Resilient Distributed Datasets (RDDs)

- Containers for operand data** passed between parallel operations
 - Read-only* (after construction) collection of data objects
 - Partitioned and distributed across workers (cluster nodes)
 - Materialized on demand from construction description
 - Can be rebuilt if a partition (data block) is lost
 - By default, cached in **main memory** – not persistent (in secondary storage) until written back



Construction of new RDDs:

- By reading in from a file e.g. in HDFS
- By partitioning and distributing a non-distributed collection (e.g., array) previously residing on one node ("*scatter*")
- By a *Map* operation: $A \rightarrow \text{List}(B)$ (elementwise transformation, filtering, ...) applied on another RDD
- Changing persistence state of a RDD:**
 - By a **caching** hint for data to be reused – if enough space in memory
 - By **materializing** (persisting, saving) to a file (and discarding its copy in memory)

data = [1, 2, 3, 4, 5]
distData = sc.parallelize(data)

cachedData = distData.cache()
distData.saveAsTextFile(...)


C. Kessler, IDA, Linköpings universitet.

6

Parallel Operations on RDDs

Spark execution model:

- **Driver program** (sequential) runs on host / master
- Operations on RDDs run on **workers**
- Collect data from workers to driver program on demand:



Parallel Collect Operations on RDDs:

■ Reduce

- Combine RDD elements using an associative binary function to produce a (scalar) result at the driver program
- Key-value pairs to reduce over are grouped by key, as in MapReduce
- Initially no shuffle&sort phase → no grouped reduction as in MapReduce

■ Collect

- Send all elements of the RDD to the driver program ("**gather**")

■ Foreach

- Pass each RDD element through a user-provided function
- Not producing another RDD (difference from Map/Filter)
- Might be used e.g. for copying data to another system

`distData = sc.parallelize(data)`
`distData.collect()`

Classification of RDD Operations

■ Transformations: Lazy, parallelizable

- Mostly variants of Map and reading from file

■ Actions: Materialization points ("push the button")

- Mostly variants of Reduce and writing back to file/master

Transformations	<code>map(f: T => U)</code>	<code>: RDD[T] => RDD[U]</code>
	<code>filter(f: T => Boolean)</code>	<code>: RDD[T] => RDD[T]</code>
	<code>flatMap(f: T => Seq[U])</code>	<code>: RDD[T] => RDD[U]</code>
	<code>sample(fraction: Float)</code>	<code>: RDD[T] => RDD[T] (Deterministic sampling)</code>
	<code>groupByKey()</code>	<code>: RDD[(K, V)] => RDD[(K, Seq[V])]</code>
	<code>reduceByKey(f: (V, V) => V)</code>	<code>: RDD[(K, V)] => RDD[(K, V)]</code>
	<code>union()</code>	<code>: (RDD[T], RDD[T]) => RDD[T]</code>
	<code>join()</code>	<code>: (RDD[(K, V)], RDD[(K, W)]) => RDD[(K, (V, W))]</code>
	<code>cogroup()</code>	<code>: (RDD[(K, V)], RDD[(K, W)]) => RDD[(K, (Seq[V], Seq[W]))]</code>
	<code>crossProduct()</code>	<code>: (RDD[T], RDD[U]) => RDD[(T, U)]</code>
	<code>mapValues(f: V => W)</code>	<code>: RDD[(K, V)] => RDD[(K, W)] (Preserves partitioning)</code>
	<code>sort(c: Comparator[K])</code>	<code>: RDD[(K, V)] => RDD[(K, V)]</code>
	<code>partitionBy(p: Partitioner[K])</code>	<code>: RDD[(K, V)] => RDD[(K, V)]</code>
	<code>count()</code>	<code>: RDD[T] => Long</code>
Actions	<code>collect()</code>	<code>: RDD[T] => Seq[T]</code>
	<code>reduce(f: (T, T) => T)</code>	<code>: RDD[T] => T</code>
	<code>lookup(k: K)</code>	<code>: RDD[(K, V)] => Seq[V] (On hash/range partitioned RDDs)</code>
	<code>save(path: String)</code>	<code>: Outputs RDD to a storage system, e.g., HDFS</code>

RDD transformations and actions available in Spark. Seq[T] denotes a sequence of elements of type T.
Table source: Zaharia et al. "Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing."

Shared Variables

■ Broadcast Variables

- Replicated shared variables – 1 copy on each worker
- Read-only for workers
- For global data needed by all workers, e.g. filtering parameters, lookup table

■ Accumulator Variables

- Residing on driver program process
- Workers can not read, only add their contributions using an associative operation
- Good for implementing counters and for global sum

Example: Text Search

■ Count lines containing errors in a large log file stored in HDFS

```
// Create a RDD from file:
file = sc.textFile("hdfs://...")

// Filter operation to create RDD containing lines with "ERROR":
errs = file.filter( lambda line: line.find("ERROR")>=0 )

// Map each line to a 1:
ones = errs.map( lambda word: (word, 1) )

// Add up the 1's using Reduce:
count = ones.reduce( lambda x, y: x+y )
```

The "lineage" of RDDs leading to the result count

- RDDs *errs* and *ones* are lazy RDDs that are never materialized to secondary storage.

- Call to **reduce** triggers computation of *ones*, which triggers computation of *errs*, which triggers reading blocks from the file.

Example: Text Search, with reuse of errs

■ Count lines containing errors in a large log file stored in HDFS

```
// Create a RDD from file:
file = sc.textFile("hdfs://...")

// Filter operation to create RDD containing lines with "ERROR":
errs = file.filter( lambda line: line.find("ERROR")>=0 )

// Cache hint that errs will be reused in another operation:
cachedErrs = errs.cache();

// Map each line to a 1:
ones = cachedErrs.map( lambda word: (word, 1) )


// Add up the 1's using Reduce:
count = ones.reduce( lambda x, y: x+y )
```



Example: Pi Calculation

■ Stochastic approximation of Pi:

- A random point (x,y) in [0,1]x[0,1] is located within quarter unit circle iff $x^2 + y^2 < 1$



def sample(p):

```
x, y = random(), random()
return 1 if x*x + y*y < 1 else 0
```

Create a RDD containing all indexes 0, ..., NUM_SAMPLES-1

```
count = sc.parallelize( xrange(0, NUM_SAMPLES) ) \
    .map(sample) \
    .reduce( lambda a, b: a + b )
```

```
print "Pi is roughly %f" % (4.0 * count / NUM_SAMPLES)
```

Example: Logistic Regression

- Iterative classification algorithm to find a hyperplane that best separates 2 sets of data points
- Gradient descent method:
 - Start at a random normal-vector (hyperplane) w
 - In each iteration, add to w an error-correction term (based on the *gradient*) that is a function of w and the data points, to improve w

```
// Read points from a text file and cache them:
points = sc.textFile(...).map(parsePoint).cache()
// Initialize w to random D-dimensional vector:
w = Vector.random(D)
// Run multiple iterations to update w:
for (i <- 1 to NUMBER_OF_ITERATIONS) {
  grad = sc.accumulator( new Vector(D) )
  for (p <- points) { // Runs in parallel:
    val s = (1/(1+exp(-p.y*(w dot p.x)))-1) * p.y
    grad += s * p.x // remotely add contribution to gradient value
  }
  w -= grad.value // correction of w
}
```


Scala pseudocode, adapted from Zaharia et al., 2010

C. Kessler

13

Spark Execution Model

- Depending on the kind of operations, the data dependencies between RDDs in the lineage graph can be **local** (elementwise) or **global** (shuffle-like)
- When a user runs an *action* on an RDD, the Spark scheduler builds a DAG of *stages* from the RDD lineage graph.
- A **stage** contains a contiguous subDAG of as many as possible operations with *local* (element-wise) dependencies between RDDs
 - The boundary of a stage is thus defined by
 - Operations with global dependencies
 - Already computed (materialized) RDD partitions.
- Execution of the operations within a stage is **pipelined**
 - intermediate results forwarded in memory
- The scheduler launches **tasks** to workers (cluster nodes) to compute missing partitions from each stage until it computes the target RDD.
- Tasks are assigned to nodes based on data locality.
 - If a task needs a partition that is available in the memory of a node, the task is sent to that node.



C. Kessler, IDA, Linköping universitet.

14

Spark Performance

Results from original paper on Spark 2010:

- Spark can outperform Hadoop by 10x in iterative machine learning jobs
- Interactive query of a 39GB data set in < 1s




Figure 2: Logistic regression performance in Hadoop and Spark.

Image source: M. Zaharia et al., 2010. © ACM

C. Kessler, IDA, Linköping universitet.

15

Using Spark

- Spark can run atop HDFS, but other implementations also exist
- Language bindings exist for Scala, Java, Python (PySpark)
 - Some minor restrictions for Python
- Spark Context object
 - The main entry point to Spark functionality
 - Represents connection to a Spark cluster
 - PySpark context `sc` is up and running from start
 - Create your own Spark context object for stand-alone applications
 - `sc = new pyspark.SparkContext(master, appName, [sparkHome], [...])`

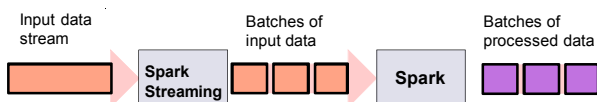
local
local[k]
spark://host:port
mesos://host:port

C. Kessler, IDA, Linköping universitet.

16

Spark Streaming

- Extension of the core Spark API for scalable, high-throughput, fault-tolerant stream processing of live data streams.
- Discretized stream** or **DStream**
 - High-level abstraction representing a continuous stream of data.
 - Internally: A continuous series of RDDs



C. Kessler, IDA, Linköping universitet.

17

Transformations on DStreams

- map(func), flatMap(func), filter(func)** – return a new DStream with **map** etc. applied to all its elements
- repartition(), union(other_stream)**
- count()** – returns a new DStream of single-element RDDs containing the number of elements in each RDD of the source DStream
- reduce(func), reduceByKey()** – aggregate each RDD of the source Dstream and return a new Dstream of single-element RDDs
- join(other_stream)** – joins 2 streams of (K,V) and (K,W) pairs to a stream of (K,(V,W)) pairs
- transform(func)** – apply arbitrary RDD-to-RDD function to each RDD in the source DStream
- ...

C. Kessler, IDA, Linköping universitet.

18

Spark Streaming Example

```
from pyspark import SparkContext
from pyspark.streaming import StreamingContext
```

```
# Create a local StreamingContext with two working threads and batch interval of 1 second:
sc = SparkContext("local[2]", "NetworkWordCount")
ssc = StreamingContext(sc, 1)
```

Run on local host, alt. cluster name

```
# Create a DStream that will connect to TCP hostname:port, like localhost:9999, as source:
lines = ssc.socketTextStream("localhost", 9999)
```

DStream of lines

```
# Split each line into words:
words = lines.flatMap( lambda line: line.split(" ") )
```

```
# Count each word in each batch:
pairs = words.map( lambda word: (word, 1) )
wordCounts = pairs.reduceByKey( lambda x, y: x + y )
```

```
# Print the first ten elements of each RDD generated in this DStream to the console:
wordCounts.pprint()
```


```
ssc.start() # Start the computation
ssc.awaitTermination() # Wait for the computation to terminate
```

C. Kessler, IDA Linköping universitet.

19

Spark Streaming: Windowing

- Can define a sliding window over a source DStream



Window length (here 3)
Slide length (here 2)
 → **Overlap size** (here 1)

Every time the window slides over a source DStream, the source RDDs that fall within the window are combined and operated upon to produce the RDDs of the windowed DStream.

Example: Reduce last 30 seconds of data, every 10 seconds:

```
windowedWordCounts = \
pairs.reduceByKeyAndWindow( lambda x, y: x + y, lambda x, y: x - y, 30, 10 )
```

C. K

20

References

- M. Zaharia, M. Chowdhury, M. Franklin, S. Shenker, I. Stoica: Spark: Cluster Computing with Working Sets. Proceedings of the 2nd USENIX conference on Hot topics in cloud computing (*HotCloud'10*), 2010, ACM.
- See also: M. Zaharia *et al.*: Apache Spark: A Unified Engine for Big Data Processing. *Communications of the ACM*, 59(11):56-65, Nov. 2016.
- Apache Spark: <http://spark.apache.org>
- A. Nandi: *Spark for Python Developers*. Packt Publishing, 2015.


C. Kessler, IDA Linköping universitet.

21

Questions for Reflection

- Why can MapReduce emulate any distributed computation?
- For a Spark program consisting of 2 subsequent Map computations, show how Spark execution differs from Hadoop/MapReduce execution.
- Given is a file containing just integer numbers. Write a Spark program that adds them up.
- Write a wordcount program for Spark.
 - Solution proposal (from spark.apache.org):


```
text_file = sc.textFile("hdfs://...")
counts = text_file.flatMap( lambda line: line.split(" ") ) \
    .map( lambda word: (word, 1) ) \
    .reduceByKey( lambda a, b: a + b )
counts.saveAsTextFile("hdfs://...")
```
 - Note – there exist many variants for formulating this.
- Modify the wordcount program by only considering words with at least 4 characters.



C. Kessler, IDA Linköping universitet.

22