

Computational Statistics (732A90) Lab4

Anubhav Dikshit(anudi287) and Thijs Quast(thiqu264)

07 February 2019

Contents

Contributions:	2
Question 1: Computations with Metropolis-Hastings	2
1. Use Metropolis-Hastings algorithm to generate samples from this distribution by using proposal distribution as log-normal $LN(Xt; 1)$, take some starting point. Plot the chain you obtained as a time series plot. What can you guess about the convergence of the chain? If there is a burn-in period, what can be the size of this period?	2
2. Perform Step 1 by using the chi-square distribution $\chi^2(\text{floor}(X(t)+1))$ as a proposal distribution, where $[x]$ is the floor function, meaning the integer part of x for positive x , i.e. $[2.95]=2$	3
3. Compare the results of Steps 1 and 2 and make conclusions.	4
4. Generate 10 MCMC sequences using the generator from Step 2 and starting points 1,2,3... or 10. Use the Gelman-Rubin method to analyze convergence of these sequences.	4
5. Estimate Integral of $xf(x)$, with limits from 0 to infinity using step 1 and 2.	5
6. The distribution generated is in fact a gamma distribution. Look in the literature and define the actual value of the integral. Compare it with the one you obtained.	6
Question 2: Gibbs sampling	6
1. Import the data to R and plot the dependence of Y on X . What kind of model is reasonable to use here?	6
2. A researcher has decided to use the following (random-walk) Bayesian model (n =number of observations), $\mu = \mu_1, \mu_2, \dots$ are unknown parameters:	7
3. Use Bayes' Theorem to get the posterior up to a constant proportionality, and then find out the distributions of $(\mu_i \mu_{-i}, Y)$, where $\sim \mu_{-i}$ is a vector containing all μ values except of μ_i	8
4. Use the distributions derived in Step 3 to implement a Gibbs sampler that uses $\mu = (0, \dots, 0)$ as a starting point. Run the Gibbs sampler to obtain 1000 values of μ and then compute the expected value of μ by using a Monte Carlo approach. Plot the expected value of μ versus X and Y versus X in the same graph. Does it seem that you have managed to remove the noise? Does it seem that the expected value of μ can catch the true underlying dependence between Y and X ?	8
5. Make a trace plot for μ_n and comment on the burn-in period and convergence.	10
Appendix	11

Contributions:

Hector Plata(hecpl268) helped us a lot in the second part of the assignment.

Question 1: Computations with Metropolis-Hastings

Consider the pdf:

$$f(X) \sim x^5 e^{-x}, x > 0$$

1. Use Metropolis-Hastings algorithm to generate samples from this distribution by using proposal distribution as log-normal $LN(X_t; 1)$, take some starting point. Plot the chain you obtained as a time series plot. What can you guess about the convergence of the chain? If there is a burn-in period, what can be the size of this period?

```
target<- function(x){
  result <- (x^5 * exp(-x))
  return(result)
}

proposed <- function(x,log_mean){
  result <- dlnorm(x,meanlog=log(log_mean), sd=1)
  return(result)
}

N <- 10^4
final <- rep(1,0,N) #Vx
x0 <- 0.004;

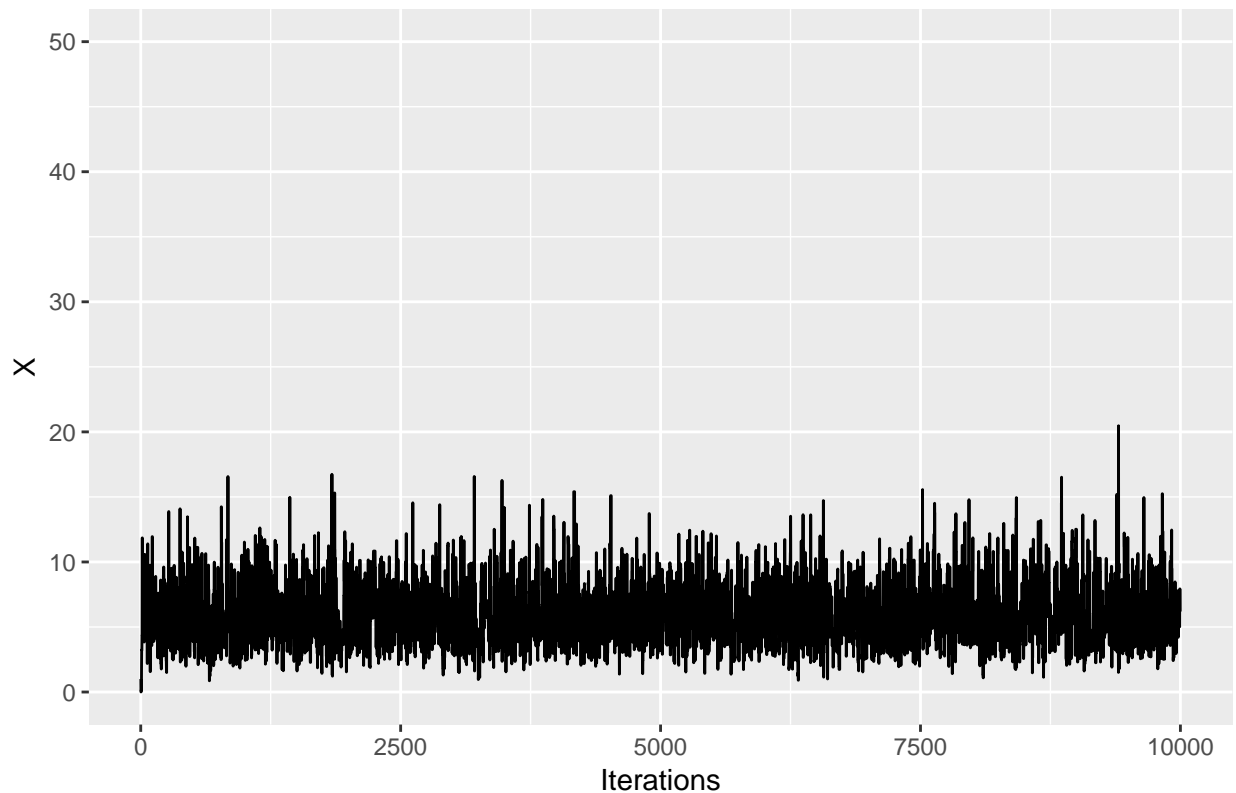
for(i in 2:N){
  xprime <- rlnorm(1, log(x0), sdlog = 1) # proposed starting value
  ratio <- min(c(1,((target(xprime)*proposed(x0, xprime))/(target(x0)*proposed(xprime, x0))))
  accept <- (runif(1) <= ratio)
  final[i] <- ifelse(accept,xprime,x0)
  x0 <- final[i]
}

step_1 <- final
mean(final)

## [1] 5.908561

ggplot() +
  geom_line(aes(x=1:N, y=final)) +
  labs(x="Iterations", y="X") +
  ylim(0, 50) +
  ggtitle("Metropolis-Hasting Sampler using Log-Normal")
```

Metropolis–Hasting Sampler using Log–Normal



Analysis: The series convergence as evident from the graph, there is a small burn-in period of about 2 iterations.

2. Perform Step 1 by using the chi-square distribution $\chi^2(\text{floor}(X(t)+1))$ as a proposal distribution, where $[x]$ is the floor function, meaning the integer part of x for positive x , i.e. $[2.95]=2$.

```
proposed_chi <- function(x){
  result <- rchisq(1, df=floor(x+1))
  return(result)
}

N <- 10^4
x0 <- 0.001 # there is a small burn in
#x0 <- 40 #starting point no burn in
final <- rep(NA,0,N)

for(i in 1:N){
  xprime <- proposed_chi(x0) # proposed starting value
  ratio <- min(c(1,((target(xprime)*proposed(x0, xprime))/(target(x0)*proposed(xprime, x0))))))
  accept <- (runif(1) < ratio)
  final[i] <- ifelse(accept,xprime,x0)
  x0 <- final[i]
}
```

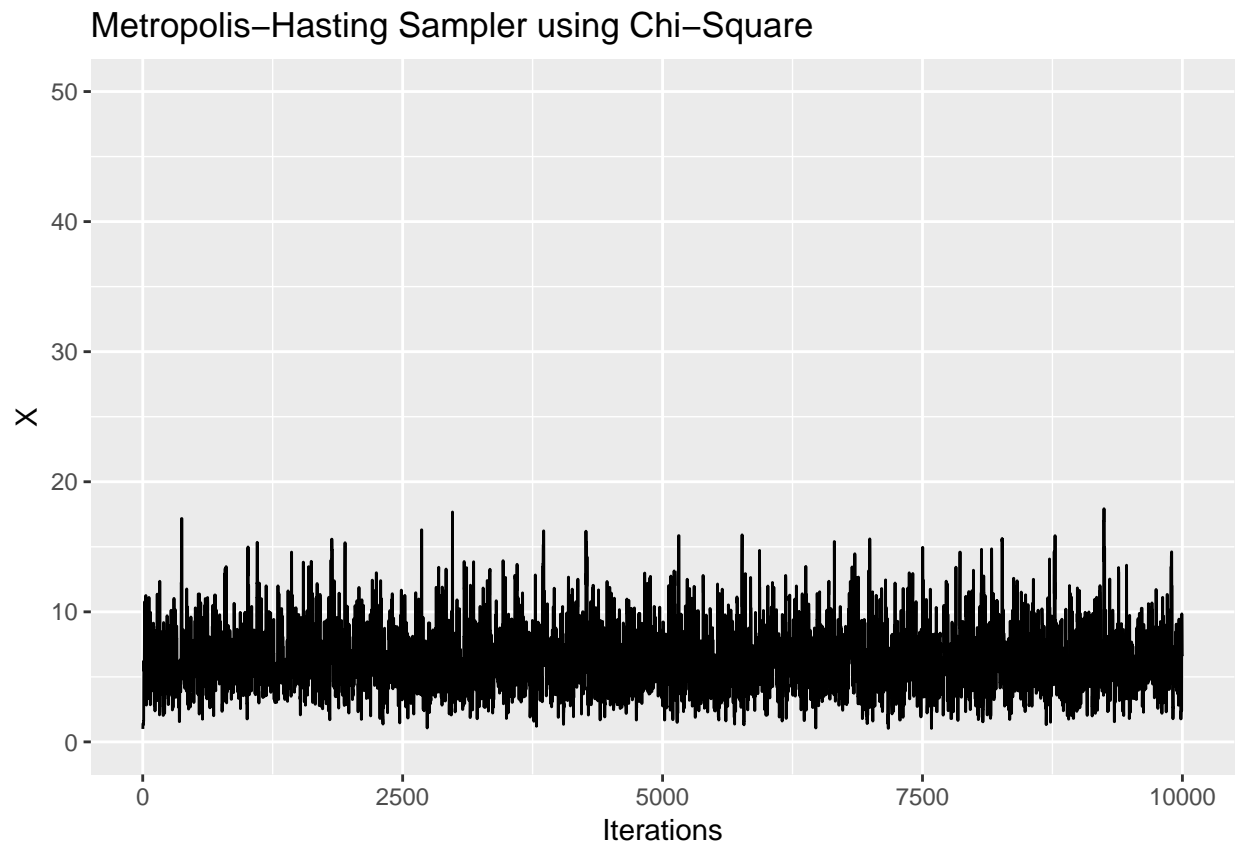
```

step_2 <- final
mean(final)

## [1] 6.234919

ggplot() +
  geom_line(aes(x=1:N, y=final)) +
  ylim(0, 50) +
  labs(x="Iterations", y="X") + ggtitle("Metropolis-Hasting Sampler using Chi-Square")

```



3. Compare the results of Steps 1 and 2 and make conclusions.

Convergence is quicker in step1 when compared to step2. Also the mean is greater when using chi-square compared to Log-normal distribution.

4. Generate 10 MCMC sequences using the generator from Step 2 and starting points 1,2,3... or 10. Use the Gelman-Rubin method to analyze convergence of these sequences.

```

target<- function(x){
  result <- (x^5 * exp(-x))
  return(result)
}

```

```

proposed_chi <- function(x){
  result <- rchisq(1, df=floor(x+1))
  return(result)
}

function_chi_fit <- function(x0){

N <- 10^4
x0 <- x0 # there is a small burn in
final <- rep(NA,0,N)

for(i in 1:N){
  xprime <- proposed_chi(x0) # proposed starting value
  ratio <- min(c(1,((target(xprime)*proposed(x0, xprime))/(target(x0)*proposed(xprime, x0)))))
  accept <- (runif(1) < ratio)
  final[i] <- ifelse(accept,xprime,x0)
  x0 <- final[i]
}
return(final)
}

all_series <- NULL
for(i in 1:10) {
  temp <- function_chi_fit(x0=i)
  temp <- coda::as.mcmc(temp)
  all_series[[i]] <- temp
}

#convergence analysis
coda::gelman.diag(all_series)

## Potential scale reduction factors:
##
##      Point est. Upper C.I.
## [1,]          1          1

```

Analysis: As we know the Upper confidence interval value should be as close to 1 as possible, this indicates that our series has converged. Our Upper C.I is 1, thus we can conclude that we series is indeed converged.

5. Estimate Integral of $xf(x)$, with limits from 0 to infinity using step 1 and 2.

When that our series convergence using Log-normal and Chi-square distribution we can use the following identity.

$$E[x] = \int xf(x)dx \approx \frac{1}{n-m} \sum_{t=m+1}^n x_t$$

```
mean(step_1)
```

```
## [1] 5.908561
```

```
mean(step_2)
```

```
## [1] 6.234919
```

6. The distribution generated is in fact a gamma distribution. Look in the literature and define the actual value of the integral. Compare it with the one you obtained.

The pdf of gamma distribution is as follows:

$$f(x|k, \theta) = \frac{1}{\Gamma(k)\theta^k} x^{k-1} e^{-\frac{x}{\theta}}$$

Analysis: Given at the beginning that gamma value constant is 120, we get $k=6$ and $\theta = 1$, thus $E[x]=k\theta = 6$. We find that chi-square function gave a much closer values and this is due to the fact that chi-square is a special case of gamma distribution.

Question 2: Gibbs sampling

A concentration of a certain chemical was measured in a water sample, and the result was stored in the data chemical.RData having the following variables: X: day of the measurement Y: measured concentration of the chemical. The instrument used to measure the concentration had certain accuracy; this is why the measurements can be treated as noisy. Your purpose is to restore the expected concentration values.

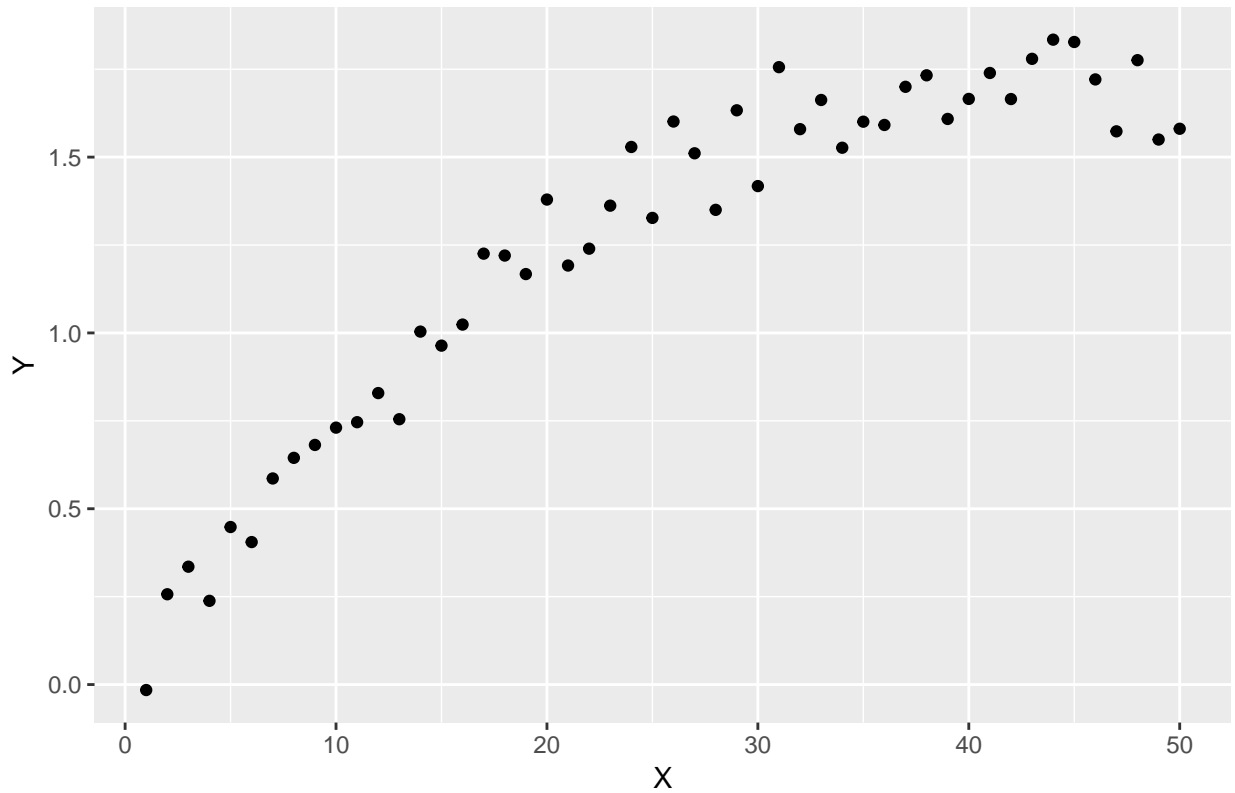
1. Import the data to R and plot the dependence of Y on X. What kind of model is reasonable to use here?

```
load(file= "chemical.Rdata")

data <- as.data.frame(cbind(X,Y))

ggplot(data=data,aes(x=X,y=Y)) + geom_point() + ggtitle("Plot of Y vs. X")
```

Plot of Y vs. X



A linear fit may be good but a 2nd order polynomial or logarithm/exponential curve might be the best fit.

2. A researcher has decided to use the following (random-walk) Bayesian model (n=number of observations), $\mu = \mu_1, \mu_2, \dots$ are unknown parameters:

$$Y \approx N(\mu_i, \text{variance} = 0.2), i = 1, 2, 3 \dots n$$

where the prior is

$$p(\mu) = 1$$

$$p(\mu_{i+1} | \mu_i) = N(\mu_i, \text{variance} = 0.2), i = 1, 2, 3 \dots n$$

Present the formula showing the likelihood $p(Y|\mu)$ and the prior $p(\mu)$.

Solution:

From chain rule we can say prior will be

$$p(\vec{\mu}) = \prod_{i=1}^{n-1} N[\mu_i, \sigma^2](\mu_i + 1)$$

Likelihood:

$$p(Y|\mu) = \prod_{i=1}^n N[\mu_i, \sigma^2](Y_i)$$

Explanation: Here we have measurements of concentration of a certain chemical over time, however there is noise in the measurements. The idea is to “restore” the trend/mean using Gibbs sampling to recreate most probable distribution of the data negating the noise due to measurements error.

3. Use Bayes’ Theorem to get the posterior up to a constant proportionality, and then find out the distributions of $(\mu_i | \mu_{-i}, Y)$, where $\sim \mu_{-i}$ is a vector containing all μ values except of μ_i

Applying Bayes theorem we get:

$$p(\mu|Y) \propto p(Y|\mu)p(\mu)$$

$$p(\mu|Y) \propto \exp\left(\frac{-1}{2\sigma^2}[(\mu_2 - \mu_1)^2 + (\mu_3 - \mu_2)^2 + \dots + (\mu_n - \mu_{n-1})^2]\right) \times \exp\left(\frac{-1}{2\sigma^2}[(Y_1 - \mu_1)^2 + (Y_2 - \mu_2)^2 + \dots + (Y_n - \mu_n)^2]\right)$$

Using property $\exp\left(\frac{-1}{d}((x-a)^2 + (x-b)^2)\right) \propto \exp\left(-\frac{[x - \frac{(a+b)}{2}]^2}{\frac{d}{2}}\right)$

Leading us to the following expression for posterior

$$p(\mu|Y) \propto \exp\left(-\frac{[\mu_n - Y_n]^2}{2\sigma^2} - \sum_{i=1}^{n-1} \frac{[\mu_i - \frac{\mu_{i+1} + Y_i}{2}]^2}{\sigma^2}\right)$$

Marginal distribution for μ_i and σ^2

$\mu_i = 1$

$$\begin{aligned} m_1 &= 2\mu_{i-1} - Y_{i-1} \\ \sigma_1 &= 2\sigma^2 \end{aligned}$$

$\mu_i = 2$

$$\begin{aligned} m_2 &= \frac{\mu_{i+1} + Y_i}{2} \\ \sigma_2 &= \sigma^2/2 \end{aligned}$$

$$p(\mu_i, \mu_{i-1} | Y) = N\left(\frac{1}{5}(2Y_i + 2\mu_{i-1} + 2\mu_{i+1} - Y_{i-1}), \frac{2\sigma^2}{5}\right)(\mu_i)$$

4. Use the distributions derived in Step 3 to implement a Gibbs sampler that uses $\mu = (0, \dots, 0)$ as a starting point. Run the Gibbs sampler to obtain 1000 values of μ and then compute the expected value of μ by using a Monte Carlo approach. Plot the expected value of μ versus X and Y versus X in the same graph. Does it seem that you have managed to remove the noise? Does it seem that the expected value of μ can catch the true underlying dependence between Y and X ?


```

gen_mu = function(i, mu, Y){

n = length(mu)
mean_mu <- rep(OL, n)
var_mu <- rep(OL, n)

mean_mu <- ifelse(i == 1, (Y[1] + mu[2])*0.5, ifelse(i == n, (2*Y[n] - Y[n-1] + 2*mu[n-1])*0.3333, (2*mu[n] + Y[n-1])*0.3333))
var_mu <- ifelse(i == 1, 0.1, ifelse(i == n, 2*0.2/5, 0.08))

# Generate mu_i using a normally distributed marginal distribution
rnorm(1, mean_mu, sqrt(var_mu))
}

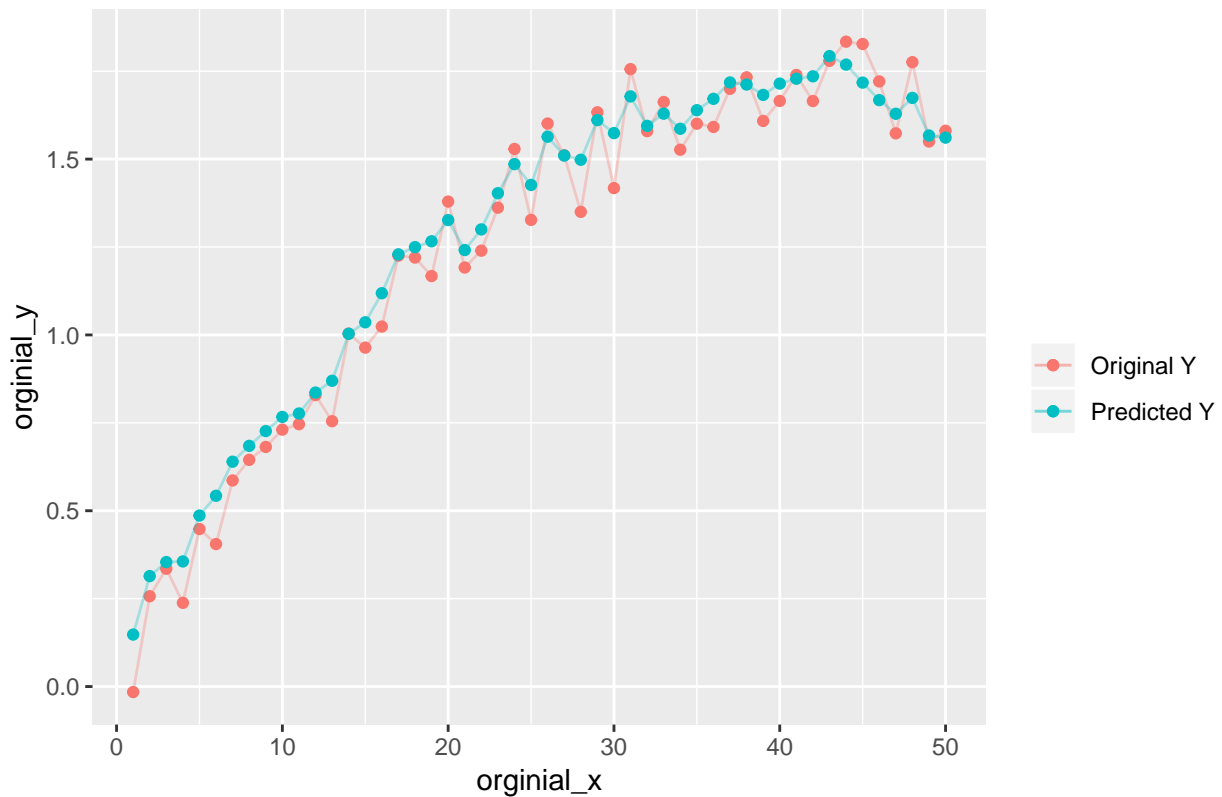
# Function to run Gibbs sampling
mcmc_gib = function(Y, num_iter){
n = length(Y)
mu0 = rep(0, length(Y))
mu = matrix(0, nrow = num_iter, ncol = n)
mu[1, ] = mu0
for (i in 2:num_iter){
latest_mu = mu[i-1, ]
for(j in 1:n){
latest_mu[j] = gen_mu(j, latest_mu, Y)
}
mu[i, ] = latest_mu
}
return(mu)
}

set.seed(12345)
result = mcmc_gib(Y, num_iter=1000)
exp_mu = colMeans(result)

ggplot(data.frame(orginial_x = X, exp_mu = exp_mu, orginial_y = Y)) +
geom_point(aes(x = orginial_x, y = orginial_y, color = "Original Y")) +
geom_line(aes(x = orginial_x, y = orginial_y, color = "Original Y"), alpha = 0.3) +
geom_point(aes(x = orginial_x, y = exp_mu, color = "Predicted Y")) +
geom_line(aes(x = orginial_x, y = exp_mu, color = "Predicted Y"), alpha = 0.3) +
labs(color="") + ggtitle("Comparison of original vs. Predicted")

```

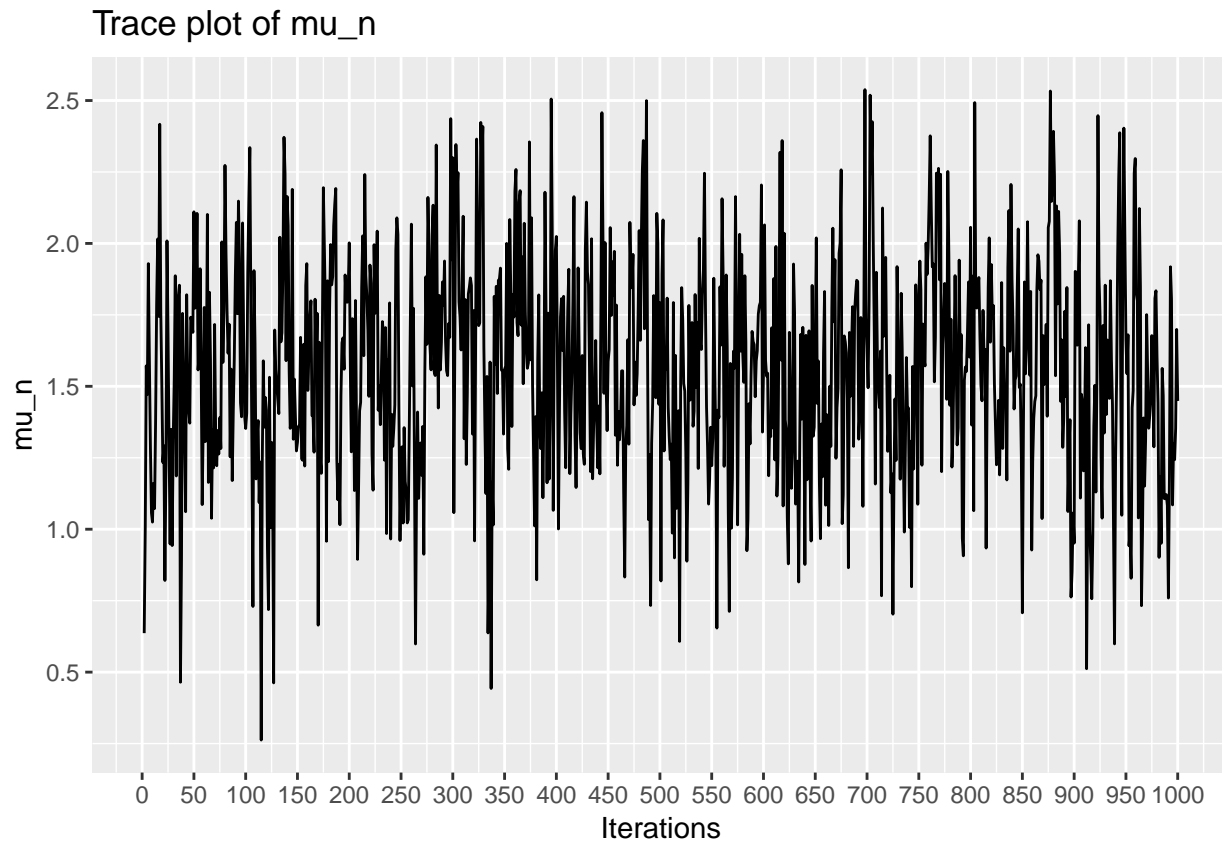
Comparison of original vs. Predicted



Analysis: As evident from the plot, our predicted value is close to the actual value.

5. Make a trace plot for μ_n and comment on the burn-in period and convergence.

```
ggplot(data.frame(x = 2:nrow(result),
y = result[2:nrow(result), ncol(result)])) +
  geom_line(aes(x = x, y = y)) +
  xlab("Iterations") + ylab("mu_n") +
  ggtitle("Trace plot of mu_n") +
  scale_x_continuous(breaks = seq(0,1000,50))
```



```
head(result[,50],10)
```

```
## [1] 0.0000000 0.6363762 1.0056683 1.5716017 1.4716282 1.9292479 1.5919426
## [8] 1.2858624 1.0570105 1.0250274
```

Analysis: As seen from values the burn-in period is about 3 iterations and convergence is indeed evident from the plot.

Appendix

```
knitr::opts_chunk$set(echo = TRUE)
options(scipen=999)
library(dplyr)
library(ggplot2)
library(coda)

target<- function(x){
  result <- (x^5 * exp(-x))
  return(result)
}

proposed <- function(x,log_mean){
  result <- dlnorm(x,meanlog=log(log_mean), sd=1)
  return(result)
}
```

```

N <- 10^4
final <- rep(1,0,N) #Vx
x0 <- 0.004;

for(i in 2:N){
  xprime <- rlnorm(1, log(x0), sdlog = 1) # proposed starting value
  ratio <- min(c(1,((target(xprime)*proposed(x0, xprime))/(target(x0)*proposed(xprime, x0))))))
  accept <- (runif(1) <= ratio)
  final[i] <- ifelse(accept,xprime,x0)
  x0 <- final[i]
}

step_1 <- final
mean(final)

ggplot() +
  geom_line(aes(x=1:N, y=final)) +
  labs(x="Iterations", y="X") +
  ylim(0, 50) +
  ggtitle("Metropolis-Hasting Sampler using Log-Normal")

proposed_chi <- function(x){
  result <- rchisq(1, df=floor(x+1))
  return(result)
}

N <- 10^4
x0 <- 0.001 # there is a small burn in
#x0 <- 40 #starting point no burn in
final <- rep(NA,0,N)

for(i in 1:N){
  xprime <- proposed_chi(x0) # proposed starting value
  ratio <- min(c(1,((target(xprime)*proposed(x0, xprime))/(target(x0)*proposed(xprime, x0))))))
  accept <- (runif(1) < ratio)
  final[i] <- ifelse(accept,xprime,x0)
  x0 <- final[i]
}

step_2 <- final
mean(final)

ggplot() +
  geom_line(aes(x=1:N, y=final)) +
  ylim(0, 50) +
  labs(x="Iterations", y="X") + ggtitle("Metropolis-Hasting Sampler using Chi-Square")

target<- function(x){
  result <- (x^5 * exp(-x))
  return(result)
}

```

```

proposed_chi <- function(x){
  result <- rchisq(1, df=floor(x+1))
  return(result)
}

function_chi_fit <- function(x0){

N <- 10^4
x0 <- x0 # there is a small burn in
final <- rep(NA,0,N)

for(i in 1:N){
  xprime <- proposed_chi(x0) # proposed starting value
  ratio <- min(c(1,((target(xprime)*proposed(x0, xprime))/(target(x0)*proposed(xprime, x0))))))
  accept <- (runif(1) < ratio)
  final[i] <- ifelse(accept,xprime,x0)
  x0 <- final[i]
}
return(final)
}

all_series <- NULL
for(i in 1:10) {
  temp <- function_chi_fit(x0=i)
  temp <- coda::as.mcmc(temp)
  all_series[[i]] <- temp
}

#convergence analysis
coda::gelman.diag(all_series)

mean(step_1)
mean(step_2)

load(file= "chemical.Rdata")

data <- as.data.frame(cbind(X,Y))

ggplot(data=data,aes(x=X,y=Y)) + geom_point() + ggtitle("Plot of Y vs. X")

gen_mu = function(i, mu, Y){

n = length(mu)
mean_mu <- rep(0L, n)
var_mu <- rep(0L, n)

mean_mu <- ifelse(i == 1, (Y[1] + mu[2])*0.5, ifelse(i == n, (2*Y[n] - Y[n-1] + 2*mu[n-1])*0.3333, (2*mu[i] + Y[n-1])*0.5))
var_mu <- ifelse(i == 1, 0.1, ifelse(i == n, 2*0.2/5, 0.08))

# Generate mu_i using a normally distributed marginal distribution
rnorm(1, mean_mu, sqrt(var_mu))

```

```

}

# Function to run Gibbs sampling
mcmc_gib = function(Y, num_iter){
  n = length(Y)
  mu0 = rep(0,length(Y))
  mu = matrix(0, nrow = num_iter, ncol = n)
  mu[1, ] = mu0
  for (i in 2:num_iter){
    latest_mu = mu[i-1, ]
    for(j in 1:n){
      latest_mu[j] = gen_mu(j, latest_mu, Y)
    }
    mu[i, ] = latest_mu
  }
  return(mu)
}

set.seed(12345)
result = mcmc_gib(Y, num_iter=1000)
exp_mu = colMeans(result)

ggplot(data.frame(orginial_x = X, exp_mu = exp_mu, orginial_y = Y)) +
  geom_point(aes(x = orginial_x, y = orginial_y, color = "Original Y")) +
  geom_line(aes(x = orginial_x, y = orginial_y, color = "Original Y"), alpha = 0.3) +
  geom_point(aes(x = orginial_x, y = exp_mu, color = "Predicted Y")) +
  geom_line(aes(x = orginial_x, y = exp_mu, color = "Predicted Y"), alpha = 0.3) +
  labs(color="") + ggtitle("Comparison of original vs. Predicted")

ggplot(data.frame(x = 2:nrow(result),
  y = result[2:nrow(result), ncol(result)])) +
  geom_line(aes(x = x, y = y)) +
  xlab("Iterations") + ylab("mu_n") +
  ggtitle("Trace plot of mu_n") +
  scale_x_continuous(breaks = seq(0,1000,50))

head(result[,50],10)

```