

machine learning(732A99) lab1 Block 2

Anubhav Dikshit(anudi287)

4 December 2018

Contents

Loading The Libraries	1
1. Your task is to evaluate the performance of Adaboost classification trees and random forests on the spam data. Specifically, provide a plot showing the error rates when the number of trees considered are 10,20,...,100. To estimate the error rates, use 2/3 of the data for training and 1/3 as hold-out test data.	2
Loading Input files	2
Splitting into Train and Test with 66% and 33% ratio.	2
Training the Model	2
2 Your task is to implement the EM algorithm for mixtures of multivariate Bernoulli distributions. Please use the template in the next page to solve the assignment. Then, use your implementation to show what happens when your mixture models has too few and too many components, i.e. set $K = 2,3,4$ and compare results. Please provide a short explanation as well.	4
Description of the EM algorithm	4
K=2	8
K=3	14
K=4	31
Analysis	57
Appendix	57

Loading The Libraries

1. Your task is to evaluate the performance of Adaboost classification trees and random forests on the spam data. Specifically, provide a plot showing the error rates when the number of trees considered are 10, 20, ..., 100. To estimate the error rates, use 2/3 of the data for training and 1/3 as hold-out test data.

Loading Input files

```
spam_data <- read.csv(file = "spambase.data", header = FALSE)
colnames(spam_data)[58] <- "Spam"
spam_data$Spam <- factor(spam_data$Spam, levels = c(0,1), labels = c("0", "1"))
```

Splitting into Train and Test with 66% and 33% ratio.

```
set.seed(12345)
n = NROW(spam_data)
id = sample(1:n, floor(n*(2/3)))
train = spam_data[id,]
test = spam_data[-id,]
```

Trainning the Model

Adaboost with varying depth

```
final_result <- NULL
for(i in seq(from = 10, to = 100, by = 10)){

  ada_model <- mboost::blackboost(Spam~.,
                                data = train,
                                family = AdaExp(),
                                control=boost_control(mstop=i))

  forest_model <- randomForest(Spam~., data = train, ntree = i)

  prediction_function <- function(model, data){
    predicted <- predict(model, newdata = data, type = c("class"))
    predict_correct <- ifelse(data$Spam == predicted, 1, 0)
    score <- sum(predict_correct)/NROW(data)
    return(score)
  }

  train_ada_model_predict <- predict(ada_model, newdata = train, type = c("class"))
  test_ada_model_predict <- predict(ada_model, newdata = test, type = c("class"))
  train_forest_model_predict <- predict(forest_model, newdata = train, type = c("class"))
  test_forest_model_predict <- predict(forest_model, newdata = test, type = c("class"))
}
```

```

test_predict_correct <- ifelse(test$Spam == test_forest_model_predict, 1, 0)
train_predict_correct <- ifelse(train$Spam == train_forest_model_predict, 1, 0)


train_ada_score <- prediction_function(ada_model, train)
test_ada_score <- prediction_function(ada_model, test)
train_forest_score <- prediction_function(forest_model, train)
test_forest_score <- prediction_function(forest_model, test)

iteration_result <- data.frame(number_of_trees = i,
                              accuracy = c(train_ada_score,
                                             test_ada_score,
                                             train_forest_score,
                                             test_forest_score),
                              type = c("train", "test", "train", "test"),
                              model = c("ADA", "ADA", "Forest", "Forest"))

final_result <- rbind(iteration_result, final_result)
}

final_result$error_rate_percentage <- 100*(1 - final_result$accuracy)
ggplot(data = final_result, aes(x = number_of_trees,
                                y = error_rate_percentage,
                                group = type, color = type)) +
  geom_point() +
  geom_line() +
  ggtitle("Error Rate vs. increase in trees") + facet_grid(rows = vars(model))

```



Analysis:

From the plots we can clearly see that ADA boosted methods uses more trees (~50) to reduce the test error, while randomforest achieves saturation in short number of trees (~10). We also see that random forest achieves less error than ADA tree for both tree and test cases.

2 Your task is to implement the EM algorithm for mixtures of multivariate Bernoulli distributions. Please use the template in the next page to solve the assignment. Then, use your implementation to show what happens when your mixture models has too few and too many components, i.e. set $K = 2, 3, 4$ and compare results. Please provide a short explanation as well.

Description of the EM algorithm

EM is an iterative expectation maximization technique. The way this works is for a given mixed distribution we guess the components of the data. This is done by first guessing the number of components and then randomly initializing the parameters of the said distribution (Mean, Variance).

Sometimes the data do not follow any known probability distribution but a mixture of known distributions such as:

$$p(x) = \sum_{k=1}^K p(k) \cdot p(x|k)$$

where $p(x|k)$ are called mixture components and $p(k)$ are called mixing coefficients: where $p(k)$ is denoted by

$$\pi_k$$

With the following conditions

$$0 \leq \pi_k \leq 1$$

and

$$\sum_k \pi_k = 1$$

We are also given that the mixture model follows a Bernoulli distribution, for Bernoulli we know that

$$\text{Bern}(x|\mu_k) = \prod_i \mu_{ki}^{x_i} (1 - \mu_{ki})^{(1-x_i)}$$

The EM algorithm for an Bernoulli mixed model is:

Set π and μ to some initial values Repeat until π and μ do not change E-step: Compute $p(z|x)$ for all k and n M-step: Set $\hat{\pi}_k$ to $\hat{\pi}_k(\text{ML})$ from likelihood estimate, do the same to μ

M step:

$$p(z_{nk}|x_n, \mu, \pi) = Z = \frac{\pi_k p(x_n|\mu_k)}{\sum_k p(x_n|\mu_k)}$$

E step:

$$\pi_k^{ML} = \frac{\sum_N p(z_{nk}|x_n, \mu, \pi)}{N}$$

$$\mu_{ki}^{ML} = \frac{\sum_n x_{ni} p(z_{nk}|x_n, \mu, \pi)}{\sum_n p(z_{nk}|x_n, \mu, \pi)}$$

The maximum likelihood of E step is:

$$\log_e p(X|\mu, \pi) = \sum_{n=1}^N \log_e \sum_{k=1}^K \pi_k \cdot p(x_n|\mu_k)$$

Code

To compare the results for $K = 2, 3, 4$, the `em_loop`-function provides a graphical analysis for every iteration. The function includes comments which explain what I did at which step to create the EM algorithm. The function will be finally run with $K = 2, 3, 4$.

```
em_loop = function(K) {
  # Initializing data
  set.seed(1234567890)
  max_it = 100 # max number of EM iterations
  min_change = 0.1 # min change in log likelihood between two consecutive EM iterations
  N = 1000 # number of training points
  D = 10 # number of dimensions
```

```

x = matrix(nrow=N, ncol = D) # training data
true_pi = vector(length = K) # true mixing coefficients
true_mu = matrix(nrow = K, ncol = D) # true conditional distributions
true_pi = c(rep(1/K, K))
if (K == 2) {
  true_mu[1,] = c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)
  true_mu[2,] = c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
  plot(true_mu[1,], type = "o", xlab = "dimension", col = "blue",
  ylim = c(0,1), main = "True")
  points(true_mu[2,], type="o", xlab = "dimension", col = "red",
  main = "True")
} else if (K == 3) {
  true_mu[1,] = c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)
  true_mu[2,] = c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
  true_mu[3,] = c(0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5)
  plot(true_mu[1,], type = "o", xlab = "dimension", col = "blue", ylim=c(0,1),
  main = "True")
  points(true_mu[2,], type = "o", xlab = "dimension", col = "red",
  main = "True")
  points(true_mu[3,], type = "o", xlab = "dimension", col = "green",
  main = "True")
} else {
  true_mu[1,] = c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)
  true_mu[2,] = c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
  true_mu[3,] = c(0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5)
  true_mu[4,] = c(0.3,0.5,0.5,0.7,0.5,0.5,0.5,0.5,0.4,0.5)
  plot(true_mu[1,], type = "o", xlab = "dimension", col = "blue",
  ylim = c(0,1), main = "True")
  points(true_mu[2,], type = "o", xlab = "dimension", col = "red",
  main = "True")
  points(true_mu[3,], type = "o", xlab = "dimension", col = "green",
  main = "True")
  points(true_mu[4,], type = "o", xlab = "dimension", col = "yellow",
  main = "True")
}

z = matrix(nrow = N, ncol = K) # fractional component assignments
pi = vector(length = K) # mixing coefficients
mu = matrix(nrow = K, ncol = D) # conditional distributions
llik = vector(length = max_it) # log likelihood of the EM iterations
# Producing the training data
for(n in 1:N) {
  k = sample(1:K, 1, prob=true_pi)
  for(d in 1:D) {
    x[n,d] = rbinom(1, 1, true_mu[k,d])
  }
}

# Random initialization of the paramters
pi = runif(K, 0.49, 0.51)
pi = pi / sum(pi)
for(k in 1:K) {
  mu[k,] = runif(D, 0.49, 0.51)
}

#EM algorithm

```

```

for(it in 1:max_it) {
  # Plotting mu
  # Defining plot title
  title = paste0("Iteration", it)
  if (K == 2) {
    plot(mu[1,], type = "o", xlab = "dimension", col = "blue", ylim = c(0,1), main = title)
    points(mu[2,], type = "o", xlab = "dimension", col = "red", main = title)
  } else if (K == 3) {
    plot(mu[1,], type = "o", xlab = "dimension", col = "blue", ylim = c(0,1), main = title)
    points(mu[2,], type = "o", xlab = "dimension", col = "red", main = title)
    points(mu[3,], type = "o", xlab = "dimension", col = "green", main = title)
  } else {
    plot(mu[1,], type = "o", xlab = "dimension", col = "blue", ylim = c(0,1), main = title)
    points(mu[2,], type = "o", xlab = "dimension", col = "red", main = title)
    points(mu[3,], type = "o", xlab = "dimension", col = "green", main = title)
    points(mu[4,], type = "o", xlab = "dimension", col = "yellow", main = title)
  }
  Sys.sleep(0.5)
  # E-step: Computation of the fractional component assignments
  for (n in 1:N) {
    # Creating empty matrix (column 1:K = p_x_given_k; column K+1 = p(x/all k)
    p_x = matrix(data = c(rep(1,K), 0), nrow = 1, ncol = K+1)
    # Calculating p(x/k) and p(x/all k)
    for (k in 1:K) {
      # Calculating p(x/k)
      for (d in 1:D) {
        p_x[1,k] = p_x[1,k] * (mu[k,d]^x[n,d]) * (1-mu[k,d])^(1-x[n,d])
      }
      p_x[1,k] = p_x[1,k] * pi[k] # weighting with pi[k]
      # Calculating p(x/all k) (denominator)
      p_x[1,K+1] = p_x[1,K+1] + p_x[1,k]
    }
    # Calculating z for n and all k
    for (k in 1:K) {
      z[n,k] = p_x[1,k] / p_x[1,K+1]
    }
  }
  # Log likelihood computation
  for (n in 1:N) {
    for (k in 1:K) {
      log_term = 0
      for (d in 1:D) {
        log_term = log_term + x[n,d] * log(mu[k,d]) + (1-x[n,d]) * log(1-mu[k,d])
      }
      llik[it] = llik[it] + z[n,k] * (log(pi[k]) + log_term)
    }
  }
  cat("iteration: ", it, "log likelihood: ", llik[it], "\n")
  flush.console()
  # Stop if the log likelihood has not changed significantly
  if (it != 1) {
    if (abs(llik[it] - llik[it-1]) < min_change) {
      break
    }
  }
}

```


```


}
}
# M-step: ML parameter estimation from the data and fractional component assignments
# Updating pi
for (k in 1:K) {
pi[k] = sum(z[,k])/N
}
# Updating mu
for (k in 1:K) {
mu[k,] = 0
for (n in 1:N) {
mu[k,] = mu[k,] + x[n,] * z[n,k]
}
mu[k,] = mu[k,] / sum(z[,k])
}
}
# Printing pi, mu and development of log likelihood at the end
return(list(
pi = pi,
mu = mu,
logLikelihoodDevelopment = plot(llik[1:it],
type = "o",
main = "Development of the log likelihood",
xlab = "iteration",
ylab = "log likelihood")
))
}

```


K=2

```
em_loop(2)
```







iteration: 1 log likelihood: -7623.897




iteration: 2 log likelihood: -7610.745




iteration: 3 log likelihood: -7463.445




iteration: 4 log likelihood: -6575.121




iteration: 5 log likelihood: -5731.559




iteration: 6 log likelihood: -5656.174




iteration: 7 log likelihood: -5648.904




iteration: 8 log likelihood: -5646.139




iteration: 9 log likelihood: -5644.608




iteration: 10 log likelihood: -5643.615




iteration: 11 log likelihood: -5642.913




iteration: 12 log likelihood: -5642.386




iteration: 13 log likelihood: -5641.977




iteration: 14 log likelihood: -5641.649




iteration: 15 log likelihood: -5641.382




iteration: 16 log likelihood: -5641.161




iteration: 17 log likelihood: -5640.975




iteration: 18 log likelihood: -5640.819




iteration: 19 log likelihood: -5640.685



iteration: 20 log likelihood: -5640.571





```
## iteration: 21 log likelihood: -5640.473
```




```
## $pi
## [1] 0.5110531 0.4889469
##
## $mu
##      [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]
## [1,] 0.4931735 0.3974606 0.5967811 0.2785480 0.6927917 0.2184957 0.8018491
## [2,] 0.4989543 0.6255823 0.3804363 0.7171478 0.3230343 0.7778699 0.2049559
##      [,8]      [,9]     [,10]
## [1,] 0.1116477 0.88054439 0.004290353
## [2,] 0.9140913 0.08997919 0.999714736
##
## $logLikelihoodDevelopment
## NULL
```

K=3


```
em_loop(3)
```




iteration: 1 log likelihood: -8029.723




iteration: 2 log likelihood: -8027.183




iteration: 3 log likelihood: -8024.696




iteration: 4 log likelihood: -8005.631




iteration: 5 log likelihood: -7877.606




iteration: 6 log likelihood: -7403.513




iteration: 7 log likelihood: -6936.919




iteration: 8 log likelihood: -6818.582




iteration: 9 log likelihood: -6791.377




iteration: 10 log likelihood: -6780.713




iteration: 11 log likelihood: -6774.958




iteration: 12 log likelihood: -6771.261




iteration: 13 log likelihood: -6768.606




iteration: 14 log likelihood: -6766.535




iteration: 15 log likelihood: -6764.815




iteration: 16 log likelihood: -6763.316




iteration: 17 log likelihood: -6761.967




iteration: 18 log likelihood: -6760.727




iteration: 19 log likelihood: -6759.572




iteration: 20 log likelihood: -6758.491




iteration: 21 log likelihood: -6757.475




iteration: 22 log likelihood: -6756.521




iteration: 23 log likelihood: -6755.625




iteration: 24 log likelihood: -6754.784




iteration: 25 log likelihood: -6753.996




iteration: 26 log likelihood: -6753.26




iteration: 27 log likelihood: -6752.571




iteration: 28 log likelihood: -6751.928




iteration: 29 log likelihood: -6751.328




iteration: 30 log likelihood: -6750.768




iteration: 31 log likelihood: -6750.246




iteration: 32 log likelihood: -6749.758




iteration: 33 log likelihood: -6749.304




iteration: 34 log likelihood: -6748.88




iteration: 35 log likelihood: -6748.484




iteration: 36 log likelihood: -6748.114




iteration: 37 log likelihood: -6747.767




iteration: 38 log likelihood: -6747.444




iteration: 39 log likelihood: -6747.14




iteration: 40 log likelihood: -6746.856




iteration: 41 log likelihood: -6746.589




iteration: 42 log likelihood: -6746.338




iteration: 43 log likelihood: -6746.102




iteration: 44 log likelihood: -6745.88




iteration: 45 log likelihood: -6745.67




iteration: 46 log likelihood: -6745.472




iteration: 47 log likelihood: -6745.285




iteration: 48 log likelihood: -6745.108




iteration: 49 log likelihood: -6744.939




iteration: 50 log likelihood: -6744.78




iteration: 51 log likelihood: -6744.627




iteration: 52 log likelihood: -6744.483




iteration: 53 log likelihood: -6744.344




iteration: 54 log likelihood: -6744.212




iteration: 55 log likelihood: -6744.086




iteration: 56 log likelihood: -6743.964




iteration: 57 log likelihood: -6743.848




iteration: 58 log likelihood: -6743.736




iteration: 59 log likelihood: -6743.628




iteration: 60 log likelihood: -6743.524



iteration: 61 log likelihood: -6743.423



iteration: 62 log likelihood: -6743.326





```
## $pi
## [1] 0.3259592 0.3044579 0.3695828
##
```


```
## $mu
##      [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]
## [1,] 0.4737193 0.3817120 0.6288021 0.3086143 0.6943731 0.1980896 0.7879447
## [2,] 0.4909874 0.4793213 0.4691560 0.4791793 0.5329895 0.4928830 0.4643990
## [3,] 0.5089571 0.5834802 0.4199272 0.7157107 0.2905703 0.7667258 0.2320784
##      [,8]      [,9]      [,10]
## [1,] 0.1349651 0.8912534 0.01937869
## [2,] 0.4902682 0.4922194 0.39798407
## [3,] 0.8516111 0.1072226 0.99981353
##
## $logLikelihoodDevelopment
## NULL
```

K=4


```
em_loop(4)
```




```
## iteration: 1 log likelihood: -8316.904
```




iteration: 2 log likelihood: -8291.114




iteration: 3 log likelihood: -8286.966




iteration: 4 log likelihood: -8264.806




iteration: 5 log likelihood: -8161.19




iteration: 6 log likelihood: -7868.89




iteration: 7 log likelihood: -7570.873




iteration: 8 log likelihood: -7445.719




iteration: 9 log likelihood: -7389.741




iteration: 10 log likelihood: -7356.803




iteration: 11 log likelihood: -7337.208




iteration: 12 log likelihood: -7326.118




iteration: 13 log likelihood: -7319.998




iteration: 14 log likelihood: -7316.6




iteration: 15 log likelihood: -7314.666




iteration: 16 log likelihood: -7313.528




iteration: 17 log likelihood: -7312.829




iteration: 18 log likelihood: -7312.367




iteration: 19 log likelihood: -7312.024




iteration: 20 log likelihood: -7311.723




iteration: 21 log likelihood: -7311.407




iteration: 22 log likelihood: -7311.036




iteration: 23 log likelihood: -7310.574




iteration: 24 log likelihood: -7309.988




iteration: 25 log likelihood: -7309.248




iteration: 26 log likelihood: -7308.322




iteration: 27 log likelihood: -7307.185




iteration: 28 log likelihood: -7305.809




iteration: 29 log likelihood: -7304.176




iteration: 30 log likelihood: -7302.273




iteration: 31 log likelihood: -7300.1




iteration: 32 log likelihood: -7297.671




iteration: 33 log likelihood: -7295.014




iteration: 34 log likelihood: -7292.171




iteration: 35 log likelihood: -7289.196




iteration: 36 log likelihood: -7286.15




iteration: 37 log likelihood: -7283.093




iteration: 38 log likelihood: -7280.079




iteration: 39 log likelihood: -7277.151




iteration: 40 log likelihood: -7274.34




iteration: 41 log likelihood: -7271.66




iteration: 42 log likelihood: -7269.116




iteration: 43 log likelihood: -7266.7




iteration: 44 log likelihood: -7264.398




iteration: 45 log likelihood: -7262.189




iteration: 46 log likelihood: -7260.051




iteration: 47 log likelihood: -7257.96




iteration: 48 log likelihood: -7255.892




iteration: 49 log likelihood: -7253.824




iteration: 50 log likelihood: -7251.733




iteration: 51 log likelihood: -7249.603




iteration: 52 log likelihood: -7247.419




iteration: 53 log likelihood: -7245.17




iteration: 54 log likelihood: -7242.853




iteration: 55 log likelihood: -7240.472




iteration: 56 log likelihood: -7238.038




iteration: 57 log likelihood: -7235.571




iteration: 58 log likelihood: -7233.095




iteration: 59 log likelihood: -7230.64




iteration: 60 log likelihood: -7228.239




iteration: 61 log likelihood: -7225.925




iteration: 62 log likelihood: -7223.725




iteration: 63 log likelihood: -7221.663




iteration: 64 log likelihood: -7219.755




iteration: 65 log likelihood: -7218.01




iteration: 66 log likelihood: -7216.431




iteration: 67 log likelihood: -7215.013




iteration: 68 log likelihood: -7213.748




iteration: 69 log likelihood: -7212.621




iteration: 70 log likelihood: -7211.62




iteration: 71 log likelihood: -7210.727




iteration: 72 log likelihood: -7209.929




iteration: 73 log likelihood: -7209.208




iteration: 74 log likelihood: -7208.552




iteration: 75 log likelihood: -7207.946




iteration: 76 log likelihood: -7207.38




iteration: 77 log likelihood: -7206.844




iteration: 78 log likelihood: -7206.327




iteration: 79 log likelihood: -7205.824




iteration: 80 log likelihood: -7205.326




iteration: 81 log likelihood: -7204.829




iteration: 82 log likelihood: -7204.327




iteration: 83 log likelihood: -7203.816




iteration: 84 log likelihood: -7203.294




iteration: 85 log likelihood: -7202.756




iteration: 86 log likelihood: -7202.201




iteration: 87 log likelihood: -7201.627




iteration: 88 log likelihood: -7201.032




iteration: 89 log likelihood: -7200.414




iteration: 90 log likelihood: -7199.773




iteration: 91 log likelihood: -7199.107




iteration: 92 log likelihood: -7198.416




iteration: 93 log likelihood: -7197.7




iteration: 94 log likelihood: -7196.957




iteration: 95 log likelihood: -7196.188




iteration: 96 log likelihood: -7195.392




iteration: 97 log likelihood: -7194.57




iteration: 98 log likelihood: -7193.722



```
## iteration: 99 log likelihood: -7192.847
```



```
## iteration: 100 log likelihood: -7191.946
```



```
## $pi
## [1] 0.2880470 0.2533761 0.2933710 0.1652060
##
## $mu
##      [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]
## [1,] 0.3714855 0.3899958 0.4790260 0.5731886 0.5022651 0.5108478 0.2835691
## [2,] 0.5199997 0.6135841 0.3891214 0.7132736 0.2722448 0.7785461 0.2168891
## [3,] 0.4383456 0.4042497 0.5489526 0.3298363 0.6578057 0.2049012 0.7825505
## [4,] 0.3428531 0.7784238 0.5591637 0.6319621 0.5167044 0.4629058 0.7311279
##      [,8]      [,9]     [,10]
## [1,] 0.3519184 0.36924863 0.48252239
## [2,] 0.9337959 0.08504806 0.99916297
## [3,] 0.1703330 0.80517853 0.04500171
## [4,] 0.6601375 0.46532151 0.48814639
##
## $logLikelihoodDevelopment
```



```
## NULL
```

Analysis

Comparing the final plots for each of the cases, it becomes clear that when the mixture model has more components ($K = 4$), the EM algorithm does not perform as accurate as for fewer components ($K = 2$ or $K = 3$). The segregation between each component gets diluted as the components get higher.

Appendix

```
knitr::opts_chunk$set(echo = TRUE)
if (!require("pacman")) install.packages("pacman")
pacman::p_load(mboost, randomForest, dplyr, ggplot2)

options(scipen = 999)

spam_data <- read.csv(file = "spambase.data", header = FALSE)
colnames(spam_data)[58] <- "Spam"
spam_data$Spam <- factor(spam_data$Spam, levels = c(0,1), labels = c("0", "1"))
set.seed(12345)
n = NROW(spam_data)
id = sample(1:n, floor(n*(2/3)))
train = spam_data[id,]
test = spam_data[-id,]

final_result <- NULL
for(i in seq(from = 10, to = 100, by = 10)){

  ada_model <- mboost::blackboost(Spam~.,
                                data = train,
                                family = AdaExp(),
                                control=boost_control(mstop=i))

  forest_model <- randomForest(Spam~., data = train, ntree = i)

  prediction_function <- function(model, data){
    predicted <- predict(model, newdata = data, type = c("class"))
    predict_correct <- ifelse(data$Spam == predicted, 1, 0)
    score <- sum(predict_correct)/NROW(data)
    return(score)
  }

  train_ada_model_predict <- predict(ada_model, newdata = train, type = c("class"))
  test_ada_model_predict <- predict(ada_model, newdata = test, type = c("class"))
  train_forest_model_predict <- predict(forest_model, newdata = train, type = c("class"))
  test_forest_model_predict <- predict(forest_model, newdata = test, type = c("class"))

  test_predict_correct <- ifelse(test$Spam == test_forest_model_predict, 1, 0)
  train_predict_correct <- ifelse(train$Spam == train_forest_model_predict, 1, 0)
```

```

train_ada_score <- prediction_function(ada_model, train)
test_ada_score <- prediction_function(ada_model, test)
train_forest_score <- prediction_function(forest_model, train)
test_forest_score <- prediction_function(forest_model, test)

iteration_result <- data.frame(number_of_trees = i,
                              accuracy = c(train_ada_score,
                                             test_ada_score,
                                             train_forest_score,
                                             test_forest_score),
                              type = c("train", "test", "train", "test"),
                              model = c("ADA", "ADA", "Forest", "Forest"))

final_result <- rbind(iteration_result, final_result)
}

final_result$error_rate_percentage <- 100*(1 - final_result$accuracy)
ggplot(data = final_result, aes(x = number_of_trees,
                                y = error_rate_percentage,
                                group = type, color = type)) +

  geom_point() +
  geom_line() +
  ggtitle("Error Rate vs. increase in trees") + facet_grid(rows = vars(model))

em_loop = function(K) {
  # Initializing data
  set.seed(1234567890)
  max_it = 100 # max number of EM iterations
  min_change = 0.1 # min change in log likelihood between two consecutive EM iterations
  N = 1000 # number of training points
  D = 10 # number of dimensions
  x = matrix(nrow=N, ncol = D) # training data
  true_pi = vector(length = K) # true mixing coefficients
  true_mu = matrix(nrow = K, ncol = D) # true conditional distributions
  true_pi = c(rep(1/K, K))
  if (K == 2) {
    true_mu[1,] = c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)
    true_mu[2,] = c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
    plot(true_mu[1,], type = "o", xlab = "dimension", col = "blue",
         ylim = c(0,1), main = "True")
    points(true_mu[2,], type="o", xlab = "dimension", col = "red",
           main = "True")
  } else if (K == 3) {
    true_mu[1,] = c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)
    true_mu[2,] = c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
    true_mu[3,] = c(0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5)
    plot(true_mu[1,], type = "o", xlab = "dimension", col = "blue", ylim=c(0,1),
         main = "True")
    points(true_mu[2,], type = "o", xlab = "dimension", col = "red",
           main = "True")
    points(true_mu[3,], type = "o", xlab = "dimension", col = "green",
           main = "True")
  }
}

```

```

main = "True")
} else {
true_mu[1,] = c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)
true_mu[2,] = c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
true_mu[3,] = c(0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5)
true_mu[4,] = c(0.3,0.5,0.5,0.7,0.5,0.5,0.5,0.5,0.4,0.5)
plot(true_mu[1,], type = "o", xlab = "dimension", col = "blue",
ylim = c(0,1), main = "True")
points(true_mu[2,], type = "o", xlab = "dimension", col = "red",
main = "True")
points(true_mu[3,], type = "o", xlab = "dimension", col = "green",
main = "True")
points(true_mu[4,], type = "o", xlab = "dimension", col = "yellow",
main = "True")
}

z = matrix(nrow = N, ncol = K) # fractional component assignments
pi = vector(length = K) # mixing coefficients
mu = matrix(nrow = K, ncol = D) # conditional distributions
llik = vector(length = max_it) # log likelihood of the EM iterations
# Producing the training data
for(n in 1:N) {
k = sample(1:K, 1, prob=true_pi)
for(d in 1:D) {
x[n,d] = rbinom(1, 1, true_mu[k,d])
}
}

# Random initialization of the paramters
pi = runif(K, 0.49, 0.51)
pi = pi / sum(pi)
for(k in 1:K) {
mu[k,] = runif(D, 0.49, 0.51)
}

#EM algorithm
for(it in 1:max_it) {
# Plotting mu
# Defining plot title
title = paste0("Iteration", it)
if (K == 2) {
plot(mu[1,], type = "o", xlab = "dimension", col = "blue", ylim = c(0,1), main = title)
points(mu[2,], type = "o", xlab = "dimension", col = "red", main = title)
} else if (K == 3) {
plot(mu[1,], type = "o", xlab = "dimension", col = "blue", ylim = c(0,1), main = title)
points(mu[2,], type = "o", xlab = "dimension", col = "red", main = title)
points(mu[3,], type = "o", xlab = "dimension", col = "green", main = title)
} else {
plot(mu[1,], type = "o", xlab = "dimension", col = "blue", ylim = c(0,1), main = title)
points(mu[2,], type = "o", xlab = "dimension", col = "red", main = title)
points(mu[3,], type = "o", xlab = "dimension", col = "green", main = title)
points(mu[4,], type = "o", xlab = "dimension", col = "yellow", main = title)
}
Sys.sleep(0.5)
# E-step: Computation of the fractional component assignments
for (n in 1:N) {

```

```

# Creating empty matrix (column 1:K = p_x_given_k; column K+1 = p(x/all k)
p_x = matrix(data = c(rep(1,K), 0), nrow = 1, ncol = K+1)
# Calculating p(x/k) and p(x/all k)
for (k in 1:K) {
  # Calculating p(x/k)
  for (d in 1:D) {
    p_x[1,k] = p_x[1,k] * (mu[k,d]^x[n,d]) * (1-mu[k,d])^(1-x[n,d])
  }
  p_x[1,k] = p_x[1,k] * pi[k] # weighting with pi[k]
  # Calculating p(x/all k) (denominator)
  p_x[1,K+1] = p_x[1,K+1] + p_x[1,k]
}
# Calculating z for n and all k
for (k in 1:K) {
  z[n,k] = p_x[1,k] / p_x[1,K+1]
}
}
# Log likelihood computation
for (n in 1:N) {
  for (k in 1:K) {
    log_term = 0
    for (d in 1:D) {
      log_term = log_term + x[n,d] * log(mu[k,d]) + (1-x[n,d]) * log(1-mu[k,d])
    }
    llik[it] = llik[it] + z[n,k] * (log(pi[k]) + log_term)
  }
}
cat("iteration: ", it, "log likelihood: ", llik[it], "\n")
flush.console()
# Stop if the log likelihood has not changed significantly
if (it != 1) {
  if (abs(llik[it] - llik[it-1]) < min_change) {
    break
  }
}
# M-step: ML parameter estimation from the data and fractional component assignments
# Updating pi
for (k in 1:K) {
  pi[k] = sum(z[,k])/N
}
# Updating mu
for (k in 1:K) {
  mu[k,] = 0
  for (n in 1:N) {
    mu[k,] = mu[k,] + x[n,] * z[n,k]
  }
  mu[k,] = mu[k,] / sum(z[,k])
}
}
# Printing pi, mu and development of log likelihood at the end
return(list(
  pi = pi,
  mu = mu,

```

```
logLikelihoodDevelopment = plot(llik[1:it],  
  type = "o",  
  main = "Development of the log likelihood",  
  xlab = "iteration",  
  ylab = "log likelihood")  
))  
}  
  
em_loop(2)  
em_loop(3)  
em_loop(4)
```