

machine learning(732A99) lab1 Block 2

Anubhav Dikshit(anudi287)

4 December 2018

Contents

| | |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------|
| Assignment 1 | 1 |
| Loading The Libraries | 1 |
| 1. Your task is to evaluate the performance of Adaboost classification trees and random forests on the spam data. Specifically, provide a plot showing the error rates when the number of trees considered are 10,20,...,100. To estimate the error rates, use 2/3 of the data for training and 1/3 as hold-out test data. | 2 |
| Loading Input files | 2 |
| Splitting into Train and Test with 66% and 33% ratio. | 2 |
| Training the Model | 2 |
| 2 Your task is to implement the EM algorithm for mixtures of multivariate Bernoulli distributions. Please use the template in the next page to solve the assignment. Then, use your implementation to show what happens when your mixture models has too few and too many components, i.e. set $K = 2,3,4$ and compare results. Please provide a short explanation as well. | 4 |
| Function for EM Algorithm | 4 |
| EM function with loops | 7 |
| $K = 2$ | 83 |
| $K = 3$ | 97 |
| $K = 4$ | 145 |
| Appendix | 209 |

Assignment 1

Loading The Libraries

1. Your task is to evaluate the performance of Adaboost classification trees and random forests on the spam data. Specifically, provide a plot showing the error rates when the number of trees considered are 10, 20, ..., 100. To estimate the error rates, use 2/3 of the data for training and 1/3 as hold-out test data.

Loading Input files

```
spam_data <- read.csv(file = "spambase.data", header = FALSE)
colnames(spam_data)[58] <- "Spam"
spam_data$Spam <- factor(spam_data$Spam, levels = c(0,1), labels = c("0", "1"))
```

Splitting into Train and Test with 66% and 33% ratio.

```
set.seed(12345)
n = NROW(spam_data)
id = sample(1:n, floor(n*(2/3)))
train = spam_data[id,]
test = spam_data[-id,]
```

Trainning the Model

Adaboost with varying depth

```
final_result <- NULL
for(i in seq(from = 10, to = 100, by = 10)){

  ada_model <- mboost::blackboost(Spam~.,
                                data = train,
                                family = AdaExp(),
                                control=boost_control(mstop=i))

  forest_model <- randomForest(Spam~., data = train, ntree = i)

  prediction_function <- function(model, data){
    predicted <- predict(model, newdata = data, type = c("class"))
    predict_correct <- ifelse(data$Spam == predicted, 1, 0)
    score <- sum(predict_correct)/NROW(data)
    return(score)
  }

  train_ada_model_predict <- predict(ada_model, newdata = train, type = c("class"))
  test_ada_model_predict <- predict(ada_model, newdata = test, type = c("class"))
  train_forest_model_predict <- predict(forest_model, newdata = train, type = c("class"))
  test_forest_model_predict <- predict(forest_model, newdata = test, type = c("class"))
}
```

```

test_predict_correct <- ifelse(test$Spam == test_forest_model_predict, 1, 0)
train_predict_correct <- ifelse(train$Spam == train_forest_model_predict, 1, 0)


train_ada_score <- prediction_function(ada_model, train)
test_ada_score <- prediction_function(ada_model, test)
train_forest_score <- prediction_function(forest_model, train)
test_forest_score <- prediction_function(forest_model, test)

iteration_result <- data.frame(number_of_trees = i,
                              accuracy = c(train_ada_score,
                                             test_ada_score,
                                             train_forest_score,
                                             test_forest_score),
                              type = c("train", "test", "train", "test"),
                              model = c("ADA", "ADA", "Forest", "Forest"))

final_result <- rbind(iteration_result, final_result)
}

final_result$error_rate_percentage <- 100*(1 - final_result$accuracy)
ggplot(data = final_result, aes(x = number_of_trees,
                               y = error_rate_percentage,
                               group = type, color = type)) +
  geom_point() +
  geom_line() +
  ggtitle("Error Rate vs. increase in trees") + facet_grid(rows = vars(model))

```



Analysis:

From the plots we can clearly see that ADA boosted methods uses more trees (~50) to reduce the test error, while randomforest achieves saturation in short number of trees (~10). We also see that random forest achieves less error than ADA tree for both tree and test cases.

2 Your task is to implement the EM algorithm for mixtures of multivariate Bernoulli distributions. Please use the template in the next page to solve the assignment. Then, use your implementation to show what happens when your mixture models has too few and too many components, i.e. set $K = 2, 3, 4$ and compare results. Please provide a short explanation as well.

Function for EM Algorithm

```
myem <- function(K){
  set.seed(1234567890)

  max_it <- 100 # max number of EM iterations
  min_change <- 0.1 # min change in log likelihood between two consecutive EM iterations
  N=1000 # number of training points
  D=10 # number of dimensions
```

```

x <- matrix(nrow=N, ncol=D) # training data
true_pi <- vector(length = K) # true mixing coefficients
true_mu <- matrix(nrow=K, ncol=D) # true conditional distributions
true_pi=c(rep(1/3, K))

if(K == 2){
  plot(true_mu[1,], type="o", col="blue", ylim=c(0,1))
  points(true_mu[2,], type="o", col="red")

  true_mu[1,]=c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)
  true_mu[2,]=c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
}else if(K == 3){
  plot(true_mu[1,], type="o", col="blue", ylim=c(0,1))
  points(true_mu[2,], type="o", col="red")
  points(true_mu[3,], type="o", col="green")

  true_mu[1,]=c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)
  true_mu[2,]=c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
  true_mu[3,]=c(0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5)
}else {
  plot(true_mu[1,], type="o", col="blue", ylim=c(0,1))
  points(true_mu[2,], type="o", col="red")
  points(true_mu[3,], type="o", col="green")
  points(true_mu[4,], type="o", col="yellow")

  true_mu[1,]=c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)
  true_mu[2,]=c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
  true_mu[3,]=c(0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5)
  true_mu[4,]=c(0.3,0.5,0.5,0.7,0.5,0.5,0.5,0.5,0.4,0.5)}

# Producing the training data
for(n in 1:N) {
  k <- sample(1:K,1,prob=true_pi)
  for(d in 1:D) {
    x[n,d] <- rbinom(1,1,true_mu[k,d])
  }
}

z <- matrix(nrow=N, ncol=K) # fractional component assignments
pi <- vector(length = K) # mixing coefficients
mu <- matrix(nrow=K, ncol=D) # conditional distributions
llik <- vector(length = max_it) # log likelihood of the EM iterations
# Random initialization of the paramters
pi <- runif(K,0.49,0.51)
pi <- pi / sum(pi)

for(k in 1:K) {
  mu[k,] <- runif(D,0.49,0.51)
}

for(it in 1:max_it) {

```

```

if(K == 2){
  plot(mu[1,], type="o", col="blue", ylim=c(0,1))
  points(mu[2,], type="o", col="red")
}else if(K == 3){
  plot(mu[1,], type="o", col="blue", ylim=c(0,1))
  points(mu[2,], type="o", col="red")
  points(mu[3,], type="o", col="green")
}else{
  plot(mu[1,], type="o", col="blue", ylim=c(0,1))
  points(mu[2,], type="o", col="red")
  points(mu[3,], type="o", col="green")
  points(mu[4,], type="o", col="yellow")}

Sys.sleep(0.5)
# E-step: Computation of the fractional component assignments

for(k in 1:K)
prod <- exp(x %*% log(t(mu))) * exp((1-x) %*% t(1-mu))

num = matrix(rep(pi,N), ncol = K, byrow = TRUE) * prod
dem = rowSums(num)
poster = num/dem

#Log likelihood computation.
llik[it] = sum(log(dem))
# Your code here
cat("iteration: ", it, "log likelihood: ", llik[it], "\n")
flush.console()
# Stop if the log likelihood has not changed significantly
if( it != 1){
  if(abs(llik[it] - llik[it-1]) < min_change){break}
}
#M-step: ML parameter estimation from the data and fractional component assignments
# Your code here
num_pi = colSums(poster)
pi = num_pi/N
mu = (t(poster) %*% x)/num_pi
}

a <- pi
b <- mu
c <- plot(llik[1:it], type="o")
result <- list(c(a,b,c))
return(result)
}

```

Analysis:

EM is an iterative expectation maximization technique. The way this works is for a given mixed distribution we guess the components of the data. This is done by first guessing the number of components and then randomly initializing the parameters of the said distribution (Mean, Variance).

Sometimes the data do not follow any known probability distribution but a mixture of known distributions such as:

$$p(x) = \sum_{k=1}^K p(k) \cdot p(x|k)$$

where $p(x|k)$ are called mixture components and $p(k)$ are called mixing coefficients: where $p(k)$ is denoted by

$$\pi_k$$

With the following conditions

$$0 \leq \pi_k \leq 1$$

and

$$\sum_k \pi_k = 1$$

$$\pi_k^{ML} = \frac{\sum_N p(z_{nk}|x_n, \mu, \pi)}{N}$$

$$\mu_{ki}^{ML} = \frac{\sum_n x_{ni} p(z_{nk}|x_n, \mu, \pi)}{\sum_n p(z_{nk}|x_n, \mu, \pi)}$$

Where

$$p(z_{nk}|x_n, \mu, \pi) = Z = \frac{\pi_k p(x_n|\mu_k)}{\sum_k p(x_n|\mu_k)}$$

EM function with loops

```
myem_loop <- function(K){
  # 2 - Mixture Models ####
  set.seed(1234567890)

  max_it <- 100 # max number of EM iterations
  min_change <- 0.1 # min change in log likelihood between two consecutive EM iterations
  N=1000 # number of training points
  D=10 # number of dimensions
  x <- matrix(nrow=N, ncol=D) # training data

  true_pi <- vector(length = K) # true mixing coefficients
  true_mu <- matrix(nrow=K, ncol=D) # true conditional distributions
  true_pi=c(rep(1/3, K))

  if (K == 2){
    true_mu[1,]=c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)
    true_mu[2,]=c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
    plot(true_mu[1,], type="o", col="blue", ylim=c(0,1))
    points(true_mu[2,], type="o", col="red")
  }else if (K == 3){
    true_mu[1,]=c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)

```

```

true_mu[2,]=c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
true_mu[3,]=c(0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5)
plot(true_mu[1,], type="o", col="blue", ylim=c(0,1))
points(true_mu[2,], type="o", col="red")
points(true_mu[3,], type="o", col="green")
}else{
true_mu[1,]=c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)
true_mu[2,]=c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
true_mu[3,]=c(0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5)
true_mu[4,]=c(0.3,0.5,0.5,0.7,0.5,0.5,0.5,0.5,0.4,0.5)
plot(true_mu[1,], type="o", col="blue", ylim=c(0,1))
points(true_mu[2,], type="o", col="red")
points(true_mu[3,], type="o", col="green")
points(true_mu[4,], type="o", col="yellow")
}

# Producing the training data
for(n in 1:N) {
  k <- sample(1:K,1,prob=true_pi)
  for(d in 1:D) {
    x[n,d] <- rbinom(1,1,true_mu[k,d])
  }
}

# number of guessed components
z <- matrix(nrow=N, ncol=K) # fractional component assignments
pi <- vector(length = K) # mixing coefficients
mu <- matrix(nrow=K, ncol=D) # conditional distributions
llik <- vector(length = max_it) # log likelihood of the EM iterations
# Random initialization of the paramters
pi <- runif(K,0.49,0.51)
pi <- pi / sum(pi)
for(k in 1:K) {
  mu[k,] <- runif(D,0.49,0.51)
}
pi
mu
for(it in 1:max_it) {
  if (K == 2){
    plot(mu[1,], type="o", col="blue", ylim=c(0,1))
    points(mu[2,], type="o", col="red")
  }else if (K == 3){
    plot(mu[1,], type="o", col="blue", ylim=c(0,1))
    points(mu[2,], type="o", col="red")
    points(mu[3,], type="o", col="green")
  }else{
    plot(mu[1,], type="o", col="blue", ylim=c(0,1))
    points(mu[2,], type="o", col="red")
    points(mu[3,], type="o", col="green")
    points(mu[4,], type="o", col="yellow")
  }
  Sys.sleep(0.5)
}

```



```
# E-step: Computation of the fractional component assignments
m <- matrix(NA, nrow = 1000, ncol = k)
```

```
#Here I create the Bernouilli probabilities, lecture 1b, slide 7. I use 3 loops to do it for the three
# not very efficient, but it works.
```

```
for (j in 1:k){
  for(each in 1:nrow(x)){
    row <- x[each,]
    vec <- c()
    for (i in 1:10) {
      a <- mu[j,i]^row[i]
      b <- a * ((1-mu[j,i])^(1-row[i]))
      vec[i] <- b
      c <- prod(vec)
    }
    m[each, j] <- c
  }
}
```

```
# Here I create a empty matrix, to store all values for the numerator of the formula on the bottom of
# slide 9, lecture 1b.
```

```
m2 <- matrix(NA, ncol = k, nrow = 1000)
```

```
# m2 stores all the values for the numerator of the formula on the bottom of slide 9, lecture 1b.
```

```
for (i in 1:1000){
  a <- pi * m[i,]
  m2[i,] <- a
}
```

```
# Sum m2 to get the denominator of the formula on the bottom of slide 9, lecture 1b.
```

```
m2_sum <- rowSums(m2)
m_final <- m2 / m2_sum
```

```
#Log likelihood computation.
```

```
ll <- matrix(nrow = 1000, ncol = K)
for (j in 1:K){
  for (i in 1:1000){
    ll[i, j] <- sum((x[i,] * log(mu[j,])) + (1 - x[i,])*log(1-mu[j,])))
  }
}
```

```
ll <- ll + pi
llnew <- m_final * ll
llik[it] <- sum(rowSums(llnew))
```

```
cat("iteration: ", it, "log likelihood: ", llik[it], "\n")
flush.console()
```

```
# Stop if the log likelihood has not changed significantly
```

```
if (it != 1){
  if (abs(llik[it] - llik[it-1]) < min_change) {break}
}
```

```
#M-step: ML parameter estimation from the data and fractional component assignments
```

```

# Create the numerator for pi, slide 9, lecture 1b.
numerator_pi <- colSums(m_final)

# Create new values for pi, stored in the vector pi_new
pi_new <- numerator_pi / N
pi_new
mnew <- matrix(NA, nrow = 1000, ncol = 10)
mu_new <- matrix(NA, nrow = K, ncol = 10)


for (j in 1:k){
  for (i in 1:1000){
    row <- x[i,] * m_final[i,j]
    mnew[i,] <- row
  }
  mnewsum <- colSums(mnew)/numerator_pi[j]
  mu_new[j,] <- mnewsum
}


# Now, to create the iterations, I have to run the code again and again, and specifying mu as new the
# created for mu. Same goes for the other variables.
mu <- mu_new
pi <- pi_new
}
z <- m_final
output1 <- pi
output2 <- mu
output3 <- plot(llik[1:it], type="o")

result <- list(c(output1, output2, output3))
return(result)
}


myem_loop(K=3)

```







iteration: 1 log likelihood: -6597.778




iteration: 2 log likelihood: -6595.239




iteration: 3 log likelihood: -6592.753




iteration: 4 log likelihood: -6573.7




iteration: 5 log likelihood: -6446.022




iteration: 6 log likelihood: -5978.865




iteration: 7 log likelihood: -5537.074




iteration: 8 log likelihood: -5429.225




iteration: 9 log likelihood: -5401.95




iteration: 10 log likelihood: -5389.023




iteration: 11 log likelihood: -5380.443




iteration: 12 log likelihood: -5373.845




iteration: 13 log likelihood: -5368.41




iteration: 14 log likelihood: -5363.759




iteration: 15 log likelihood: -5359.682




iteration: 16 log likelihood: -5356.051




iteration: 17 log likelihood: -5352.782




iteration: 18 log likelihood: -5349.816




iteration: 19 log likelihood: -5347.113




iteration: 20 log likelihood: -5344.641




iteration: 21 log likelihood: -5342.375




iteration: 22 log likelihood: -5340.295




iteration: 23 log likelihood: -5338.385




iteration: 24 log likelihood: -5336.63




iteration: 25 log likelihood: -5335.015




iteration: 26 log likelihood: -5333.529




iteration: 27 log likelihood: -5332.16




iteration: 28 log likelihood: -5330.9




iteration: 29 log likelihood: -5329.738




iteration: 30 log likelihood: -5328.666




iteration: 31 log likelihood: -5327.676




iteration: 32 log likelihood: -5326.762




iteration: 33 log likelihood: -5325.917




iteration: 34 log likelihood: -5325.135




iteration: 35 log likelihood: -5324.41




iteration: 36 log likelihood: -5323.739




iteration: 37 log likelihood: -5323.115




iteration: 38 log likelihood: -5322.537




iteration: 39 log likelihood: -5321.999




iteration: 40 log likelihood: -5321.498




iteration: 41 log likelihood: -5321.031




iteration: 42 log likelihood: -5320.596




iteration: 43 log likelihood: -5320.19




iteration: 44 log likelihood: -5319.81




iteration: 45 log likelihood: -5319.454




iteration: 46 log likelihood: -5319.121




iteration: 47 log likelihood: -5318.809




iteration: 48 log likelihood: -5318.515




iteration: 49 log likelihood: -5318.239




iteration: 50 log likelihood: -5317.979




iteration: 51 log likelihood: -5317.734




iteration: 52 log likelihood: -5317.503




iteration: 53 log likelihood: -5317.284




iteration: 54 log likelihood: -5317.077




iteration: 55 log likelihood: -5316.881




iteration: 56 log likelihood: -5316.695




iteration: 57 log likelihood: -5316.518




iteration: 58 log likelihood: -5316.349




iteration: 59 log likelihood: -5316.189




iteration: 60 log likelihood: -5316.036




iteration: 61 log likelihood: -5315.89




iteration: 62 log likelihood: -5315.75




iteration: 63 log likelihood: -5315.616




iteration: 64 log likelihood: -5315.487




iteration: 65 log likelihood: -5315.364




iteration: 66 log likelihood: -5315.246




iteration: 67 log likelihood: -5315.132




iteration: 68 log likelihood: -5315.022




iteration: 69 log likelihood: -5314.916



iteration: 70 log likelihood: -5314.814




iteration: 71 log likelihood: -5314.715




```
## [[1]]
## [1] 0.32411556 0.30717327 0.36871116 0.47383554 0.49067297 0.50907301
## [7] 0.38106095 0.47962547 0.58355730 0.62860603 0.47050854 0.41965465
## [13] 0.30784064 0.47936473 0.71594258 0.69512193 0.53269167 0.29039403
## [19] 0.19764738 0.49158731 0.76736739 0.78836950 0.46597001 0.23146477
## [25] 0.13424145 0.48919757 0.85221677 0.89207742 0.49348782 0.10652662
## [31] 0.01757154 0.39918885 0.99992807
```


K = 2

```
myem(K=2)
```







iteration: 1 log likelihood: -954.7133




iteration: 2 log likelihood: -957.1002




iteration: 3 log likelihood: -944.9229




iteration: 4 log likelihood: -857.3443




iteration: 5 log likelihood: -464.0063




iteration: 6 log likelihood: 50.2616




iteration: 7 log likelihood: 177.0235




iteration: 8 log likelihood: 189.1059




iteration: 9 log likelihood: 190.0362




iteration: 10 log likelihood: 189.9033



iteration: 11 log likelihood: 189.7476




iteration: 12 log likelihood: 189.6572




```
## [[1]]
## [1] 0.48293133107 0.51706866893 0.49007966733 0.50152946699 0.39332820172
## [6] 0.61703504228 0.59792428535 0.39113495199 0.27182539774 0.69957245640
## [11] 0.70178609664 0.33474375189 0.19877624614 0.76586485831 0.81088543925
## [16] 0.22897928765 0.08857255800 0.89200054923 0.90675481733 0.10849562636
## [21] 0.00008952955 0.94950012035
```


K = 3

```
myem(K=3)
```







iteration: 1 log likelihood: -912.7567




iteration: 2 log likelihood: -932.1921




iteration: 3 log likelihood: -932.0234




iteration: 4 log likelihood: -931.2587




iteration: 5 log likelihood: -927.8881




iteration: 6 log likelihood: -913.454




iteration: 7 log likelihood: -858.0583




iteration: 8 log likelihood: -709.6665




iteration: 9 log likelihood: -524.1097




iteration: 10 log likelihood: -433.1614




iteration: 11 log likelihood: -409.3331




iteration: 12 log likelihood: -405.2132




iteration: 13 log likelihood: -405.7233




iteration: 14 log likelihood: -407.1621




iteration: 15 log likelihood: -408.6475




iteration: 16 log likelihood: -409.9879




iteration: 17 log likelihood: -411.1645




iteration: 18 log likelihood: -412.1979




iteration: 19 log likelihood: -413.1139




iteration: 20 log likelihood: -413.934




iteration: 21 log likelihood: -414.675




iteration: 22 log likelihood: -415.3492




iteration: 23 log likelihood: -415.9659




iteration: 24 log likelihood: -416.532




iteration: 25 log likelihood: -417.0528




iteration: 26 log likelihood: -417.5328




iteration: 27 log likelihood: -417.9753




iteration: 28 log likelihood: -418.3836




iteration: 29 log likelihood: -418.7601




iteration: 30 log likelihood: -419.1074




iteration: 31 log likelihood: -419.4277




iteration: 32 log likelihood: -419.7229




iteration: 33 log likelihood: -419.995




iteration: 34 log likelihood: -420.2457




iteration: 35 log likelihood: -420.4767




iteration: 36 log likelihood: -420.6895




iteration: 37 log likelihood: -420.8856




iteration: 38 log likelihood: -421.0663




iteration: 39 log likelihood: -421.2329




iteration: 40 log likelihood: -421.3865




iteration: 41 log likelihood: -421.5282




iteration: 42 log likelihood: -421.659




iteration: 43 log likelihood: -421.7797




iteration: 44 log likelihood: -421.8913



iteration: 45 log likelihood: -421.9945




iteration: 46 log likelihood: -422.09




```
## [[1]]
## [1] 0.1679716522 0.2034249451 0.6286034026 0.4808697447 0.4735292886
## [6] 0.5009515437 0.3505033441 0.4223594793 0.5428015802 0.6091318393
## [11] 0.6067582335 0.4410624521 0.3556548071 0.2768901990 0.6282716562
## [16] 0.7016525159 0.6775124264 0.3823067747 0.1914725060 0.2364291675
## [21] 0.6645564670 0.7465112147 0.7631736115 0.3235087926 0.0563854887
## [26] 0.2226418271 0.7210236740 0.8485479479 0.8448194739 0.2634580777
## [31] 0.0005534402 0.0190935069 0.7843147843
```


K = 4

```
myem(K=4)
```







iteration: 1 log likelihood: -800.5436




iteration: 2 log likelihood: -842.949




iteration: 3 log likelihood: -842.6806




iteration: 4 log likelihood: -841.7499




iteration: 5 log likelihood: -838.7414




iteration: 6 log likelihood: -829.4624




iteration: 7 log likelihood: -803.3592




iteration: 8 log likelihood: -744.3623




iteration: 9 log likelihood: -658.0191




iteration: 10 log likelihood: -588.2999




iteration: 11 log likelihood: -553.5615




iteration: 12 log likelihood: -538.8823




iteration: 13 log likelihood: -531.9182




iteration: 14 log likelihood: -527.7567




iteration: 15 log likelihood: -524.8526




iteration: 16 log likelihood: -522.7751




iteration: 17 log likelihood: -521.3929




iteration: 18 log likelihood: -520.6263




iteration: 19 log likelihood: -520.391




iteration: 20 log likelihood: -520.5983




iteration: 21 log likelihood: -521.1652




iteration: 22 log likelihood: -522.0204




iteration: 23 log likelihood: -523.1059




iteration: 24 log likelihood: -524.3754




iteration: 25 log likelihood: -525.7912




iteration: 26 log likelihood: -527.3207




iteration: 27 log likelihood: -528.9346




iteration: 28 log likelihood: -530.6046




iteration: 29 log likelihood: -532.304




iteration: 30 log likelihood: -534.0069




iteration: 31 log likelihood: -535.6895




iteration: 32 log likelihood: -537.3305




iteration: 33 log likelihood: -538.912




iteration: 34 log likelihood: -540.4198




iteration: 35 log likelihood: -541.8433




iteration: 36 log likelihood: -543.1756




iteration: 37 log likelihood: -544.4133




iteration: 38 log likelihood: -545.5555




iteration: 39 log likelihood: -546.6036




iteration: 40 log likelihood: -547.5609




iteration: 41 log likelihood: -548.4315




iteration: 42 log likelihood: -549.2208




iteration: 43 log likelihood: -549.9344




iteration: 44 log likelihood: -550.5781




iteration: 45 log likelihood: -551.1577




iteration: 46 log likelihood: -551.6789




iteration: 47 log likelihood: -552.1471




iteration: 48 log likelihood: -552.5674




iteration: 49 log likelihood: -552.9443




iteration: 50 log likelihood: -553.2824




iteration: 51 log likelihood: -553.5855




iteration: 52 log likelihood: -553.8573




iteration: 53 log likelihood: -554.101




iteration: 54 log likelihood: -554.3194




iteration: 55 log likelihood: -554.5153




iteration: 56 log likelihood: -554.691




iteration: 57 log likelihood: -554.8485




iteration: 58 log likelihood: -554.9898




iteration: 59 log likelihood: -555.1165




iteration: 60 log likelihood: -555.2301



iteration: 61 log likelihood: -555.3319



iteration: 62 log likelihood: -555.4231



```
## [[1]]
## [1] 0.0681207107 0.7239375809 0.1144285078 0.0935132006 0.3956837783
## [6] 0.4293539072 0.4323432998 0.3929703185 0.4162506305 0.5547107213
## [11] 0.4023142592 0.4174015318 0.5420279926 0.4599340095 0.6093481972
## [16] 0.5388153590 0.3444983068 0.6261407708 0.3315033116 0.3455369524
## [21] 0.6696117509 0.4227075553 0.6799271777 0.6690887905 0.2251983009
## [26] 0.5941304803 0.1244291484 0.2278577477 0.7389032359 0.3920562642
## [31] 0.7312641782 0.7396138061 0.1989570475 0.6376090927 0.0220675386
## [36] 0.2067362122 0.7733978028 0.3148516085 0.7696255465 0.7733194550
## [41] 0.0036278205 0.6703401502 0.0000374818 0.0049635197
```

Appendix

```
knitr::opts_chunk$set(echo = TRUE)
if (!require("pacman")) install.packages("pacman")
pacman::p_load(mboost, randomForest, dplyr, ggplot2)

options(scipen = 999)

spam_data <- read.csv(file = "spambase.data", header = FALSE)
colnames(spam_data)[58] <- "Spam"
spam_data$Spam <- factor(spam_data$Spam, levels = c(0,1), labels = c("0", "1"))
set.seed(12345)
n = NROW(spam_data)
```

```

id = sample(1:n, floor(n*(2/3)))
train = spam_data[id,]
test = spam_data[-id,]

final_result <- NULL
for(i in seq(from = 10, to = 100, by = 10)){

ada_model <- mboost::blackboost(Spam~.,
                                data = train,
                                family = AdaExp(),
                                control=boost_control(mstop=i))

forest_model <- randomForest(Spam~., data = train, ntree = i)

prediction_function <- function(model, data){
  predicted <- predict(model, newdata = data, type = c("class"))
  predict_correct <- ifelse(data$Spam == predicted, 1, 0)
  score <- sum(predict_correct)/NROW(data)
  return(score)
}

train_ada_model_predict <- predict(ada_model, newdata = train, type = c("class"))
test_ada_model_predict <- predict(ada_model, newdata = test, type = c("class"))
train_forest_model_predict <- predict(forest_model, newdata = train, type = c("class"))
test_forest_model_predict <- predict(forest_model, newdata = test, type = c("class"))

test_predict_correct <- ifelse(test$Spam == test_forest_model_predict, 1, 0)
train_predict_correct <- ifelse(train$Spam == train_forest_model_predict, 1, 0)

train_ada_score <- prediction_function(ada_model, train)
test_ada_score <- prediction_function(ada_model, test)
train_forest_score <- prediction_function(forest_model, train)
test_forest_score <- prediction_function(forest_model, test)

iteration_result <- data.frame(number_of_trees = i,
                              accuracy = c(train_ada_score,
                                             test_ada_score,
                                             train_forest_score,
                                             test_forest_score),
                              type = c("train", "test", "train", "test"),
                              model = c("ADA", "ADA", "Forest", "Forest"))

final_result <- rbind(iteration_result, final_result)
}

final_result$error_rate_percentage <- 100*(1 - final_result$accuracy)
ggplot(data = final_result, aes(x = number_of_trees,
                                y = error_rate_percentage,
                                group = type, color = type)) +

```

```

geom_point() +
geom_line() +
ggtitle("Error Rate vs. increase in trees") + facet_grid(rows = vars(model))

myem <- function(K){
  set.seed(1234567890)

  max_it <- 100 # max number of EM iterations
  min_change <- 0.1 # min change in log likelihood between two consecutive EM iterations
  N=1000 # number of training points
  D=10 # number of dimensions
  x <- matrix(nrow=N, ncol=D) # training data
  true_pi <- vector(length = K) # true mixing coefficients
  true_mu <- matrix(nrow=K, ncol=D) # true conditional distributions
  true_pi=c(rep(1/3, K))

  if(K == 2){
    plot(true_mu[1,], type="o", col="blue", ylim=c(0,1))
    points(true_mu[2,], type="o", col="red")

    true_mu[1,]=c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)
    true_mu[2,]=c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
  }else if(K == 3){
    plot(true_mu[1,], type="o", col="blue", ylim=c(0,1))
    points(true_mu[2,], type="o", col="red")
    points(true_mu[3,], type="o", col="green")

    true_mu[1,]=c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)
    true_mu[2,]=c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
    true_mu[3,]=c(0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5)
  }else {
    plot(true_mu[1,], type="o", col="blue", ylim=c(0,1))
    points(true_mu[2,], type="o", col="red")
    points(true_mu[3,], type="o", col="green")
    points(true_mu[4,], type="o", col="yellow")

    true_mu[1,]=c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)
    true_mu[2,]=c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
    true_mu[3,]=c(0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5)
    true_mu[4,]=c(0.3,0.5,0.5,0.7,0.5,0.5,0.5,0.5,0.4,0.5)}

  # Producing the training data
  for(n in 1:N) {
    k <- sample(1:K,1,prob=true_pi)
    for(d in 1:D) {
      x[n,d] <- rbinom(1,1,true_mu[k,d])
    }
  }

  z <- matrix(nrow=N, ncol=K) # fractional component assignments
  pi <- vector(length = K) # mixing coefficients
  mu <- matrix(nrow=K, ncol=D) # conditional distributions
  llik <- vector(length = max_it) # log likelihood of the EM iterations

```

```

# Random initialization of the paramters
pi <- runif(K,0.49,0.51)
pi <- pi / sum(pi)

for(k in 1:K) {
mu[k,] <- runif(D,0.49,0.51)
}

for(it in 1:max_it) {

if(K == 2){
  plot(mu[1,], type="o", col="blue", ylim=c(0,1))
  points(mu[2,], type="o", col="red")
}else if(K == 3){
  plot(mu[1,], type="o", col="blue", ylim=c(0,1))
  points(mu[2,], type="o", col="red")
  points(mu[3,], type="o", col="green")
}else{
  plot(mu[1,], type="o", col="blue", ylim=c(0,1))
  points(mu[2,], type="o", col="red")
  points(mu[3,], type="o", col="green")
  points(mu[4,], type="o", col="yellow")}

Sys.sleep(0.5)
# E-step: Computation of the fractional component assignments

for(k in 1:K)
prod <- exp(x %*% log(t(mu))) * exp((1-x) %*% t(1-mu))

num = matrix(rep(pi,N), ncol = K, byrow = TRUE) * prod
dem = rowSums(num)
poster = num/dem

#Log likelihood computation.
llik[it] = sum(log(dem))
# Your code here
cat("iteration: ", it, "log likelihood: ", llik[it], "\n")
flush.console()
# Stop if the lok likelihood has not changed significantly
if( it != 1){
if(abs(llik[it] - llik[it-1]) < min_change){break}
}

#M-step: ML parameter estimation from the data and fractional component assignments
# Your code here
num_pi = colSums(poster)
pi = num_pi/N
mu = (t(poster) %*% x)/num_pi
}

```

```

a <- pi
b <- mu
c <- plot(lik[1:it], type="o")
result <- list(c(a,b,c))
return(result)
}
myem_loop <- function(K){
  # 2 - Mixture Models ####
  set.seed(1234567890)

  max_it <- 100 # max number of EM iterations
  min_change <- 0.1 # min change in log likelihood between two consecutive EM iterations
  N=1000 # number of training points
  D=10 # number of dimensions
  x <- matrix(nrow=N, ncol=D) # training data

  true_pi <- vector(length = K) # true mixing coefficients
  true_mu <- matrix(nrow=K, ncol=D) # true conditional distributions
  true_pi=c(rep(1/3, K))

  if (K == 2){
    true_mu[1,]=c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)
    true_mu[2,]=c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
    plot(true_mu[1,], type="o", col="blue", ylim=c(0,1))
    points(true_mu[2,], type="o", col="red")
  }else if (K == 3){
    true_mu[1,]=c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)
    true_mu[2,]=c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
    true_mu[3,]=c(0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5)
    plot(true_mu[1,], type="o", col="blue", ylim=c(0,1))
    points(true_mu[2,], type="o", col="red")
    points(true_mu[3,], type="o", col="green")
  }else{
    true_mu[1,]=c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)
    true_mu[2,]=c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
    true_mu[3,]=c(0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5)
    true_mu[4,]=c(0.3,0.5,0.5,0.7,0.5,0.5,0.5,0.5,0.4,0.5)
    plot(true_mu[1,], type="o", col="blue", ylim=c(0,1))
    points(true_mu[2,], type="o", col="red")
    points(true_mu[3,], type="o", col="green")
    points(true_mu[4,], type="o", col="yellow")
  }

  # Producing the training data
  for(n in 1:N) {
    k <- sample(1:K,1,prob=true_pi)
    for(d in 1:D) {
      x[n,d] <- rbinom(1,1,true_mu[k,d])
    }
  }
}

```

```

# number of guessed components
z <- matrix(nrow=N, ncol=K) # fractional component assignments
pi <- vector(length = K) # mixing coefficients
mu <- matrix(nrow=K, ncol=D) # conditional distributions
llik <- vector(length = max_it) # log likelihood of the EM iterations
# Random initialization of the parameters
pi <- runif(K,0.49,0.51)
pi <- pi / sum(pi)
for(k in 1:K) {
  mu[k,] <- runif(D,0.49,0.51)
}
pi
mu
for(it in 1:max_it) {
  if (K == 2){
    plot(mu[1,], type="o", col="blue", ylim=c(0,1))
    points(mu[2,], type="o", col="red")
  }else if (K == 3){
    plot(mu[1,], type="o", col="blue", ylim=c(0,1))
    points(mu[2,], type="o", col="red")
    points(mu[3,], type="o", col="green")
  }else{
    plot(mu[1,], type="o", col="blue", ylim=c(0,1))
    points(mu[2,], type="o", col="red")
    points(mu[3,], type="o", col="green")
    points(mu[4,], type="o", col="yellow")
  }
  Sys.sleep(0.5)
  # E-step: Computation of the fractional component assignments
  m <- matrix(NA, nrow = 1000, ncol = k)

  #Here I create the Bernoulli probabilities, lecture 1b, slide 7. I use 3 loops to do it for the three
  # not very efficient, but it works.
  for (j in 1:k){
    for(each in 1:nrow(x)){
      row <- x[each,]
      vec <- c()
      for (i in 1:10) {
        a <- mu[j,i]^row[i]
        b <- a * ((1-mu[j,i])^(1-row[i]))
        vec[i] <- b
        c <- prod(vec)
      }
      m[each, j] <- c
    }
  }

  # Here I create a empty matrix, to store all values for the numerator of the formula on the bottom of
  # slide 9, lecture 1b.
  m2 <- matrix(NA, ncol = k, nrow = 1000)

  # m2 stores all the values for the numerator of the formula on the bottom of slide 9, lecture 1b.
  for (i in 1:1000){

```

```

    a <- pi * m[i,]
    m2[i,] <- a
  }

  # Sum m2 to get the denominator of the formula on the bottom of slide 9, lecture 1b.
  m2_sum <- rowSums(m2)
  m_final <- m2 / m2_sum

  #Log likelihood computation.
  ll <- matrix(nrow = 1000, ncol = K)
  for (j in 1:K){
    for (i in 1:1000){
      ll[i, j] <- sum(((x[i,] * log(mu[j,])) + (1 - x[i,])*log(1-mu[j,])))
    }
  }

  ll <- ll + pi
  llnew <- m_final * ll
  llik[it] <- sum(rowSums(llnew))

  cat("iteration: ", it, "log likelihood: ", llik[it], "\n")
  flush.console()
  # Stop if the log likelihood has not changed significantly
  if (it != 1){
    if (abs(llik[it] - llik[it-1]) < min_change) {break}
  }
  #M-step: ML parameter estimation from the data and fractional component assignments

  # Create the numerator for pi, slide 9, lecture 1b.
  numerator_pi <- colSums(m_final)

  # Create new values for pi, stored in the vector pi_new
  pi_new <- numerator_pi / N
  pi_new
  mnew <- matrix(NA, nrow = 1000, ncol = 10)
  mu_new <- matrix(NA, nrow = K, ncol = 10)

  for (j in 1:k){
    for (i in 1:1000){
      row <- x[i,] * m_final[i,j]
      mnew[i,] <- row
    }
    mnewsum <- colSums(mnew)/numerator_pi[j]
    mu_new[j,] <- mnewsum
  }

  # Now, to create the iterations, I have to run the code again and again, and specifying mu as new the
  # created for mu. Same goes for the other variables.
  mu <- mu_new
  pi <- pi_new
}
z <- m_final
output1 <- pi

```

```
output2 <- mu
output3 <- plot(llik[1:it], type="o")

result <- list(c(output1, output2, output3))
return(result)
}
myem_loop(K=3)

myem(K=2)
myem(K=3)
myem(K=4)
```