

lab1__block2

Thijs Quast

30-11-2018

Contents

Question 1 Ensemble Methods	2
Question 2 Mixture Models	4
Appendix	96

Question 1 Ensemble Methods

```
# Loading packages and importing files ####
library(mboost)

## Loading required package: parallel
## Loading required package: stabs
## This is mboost 2.9-1. See 'package?mboost' and 'news(package = "mboost")'
## for a complete list of changes.
library(randomForest)

## randomForest 4.6-14
## Type rfNews() to see new features/changes/bug fixes.
library(ggplot2)

##
## Attaching package: 'ggplot2'
## The following object is masked from 'package:randomForest':
##
##     margin
## The following object is masked from 'package:mboost':
##
##     %+%

sp <- read.csv2("spambase.csv", header = FALSE, sep = ",", stringsAsFactors = FALSE)
num_sp <- data.frame(data.matrix(sp))
num_sp$V58 <- factor(num_sp$V58)

# shuffling data and dividing into train and test ####
n <- dim(num_sp)[1]
ncol <- dim(num_sp)[2]
set.seed(1234567890)
id <- sample(1:n, floor(n*(2/3)))
train <- num_sp[id,]
test <- num_sp[-id,]

# Adaboost
ntree <- c(10, 20, 30, 40, 50, 60, 70, 80, 90, 100)
error <- c()

for (i in seq(from = 10, to = 100, by = 10)){
  bb <- blackboost(V58 ~., data = train, control = boost_control(mstop = i), family = AdaExp())
  bb_predict <- predict(bb, newdata = test, type = c("class"))
  confusion_bb <- table(test$V58, bb_predict)
  miss_class_bb <- (confusion_bb[1,2] + confusion_bb[2,1])/nrow(test)
  error[(i/10)] <- miss_class_bb
}

error_df <- data.frame(cbind(ntree, error))

# Random forest ####
ntree_rf <- c(10, 20, 30, 40, 50, 60, 70, 80, 90, 100)
```

```

error_rf <- c()

for (i in seq(from = 10, to = 100, by = 10)){
  rf <- randomForest(V58 ~., data = train, ntree= 10)
  rf_predict <- predict(rf, newdata = test, type = c("class"))
  confusion_rf <- table(test$V58, rf_predict)
  miss_class_rf <- (confusion_rf[1,2] + confusion_rf[2,1])/nrow(test)
  error_rf[i/10] <- miss_class_rf
}


error_df_rf <- data.frame(cbind(ntree_rf, error_rf))

df <- cbind(error_df, error_df_rf)
df <- df[, -3]

plot_final <- ggplot(df, aes(ntree)) +
  geom_line(aes(y=error, color = "Adaboost")) +
  geom_line(aes(y=error_rf, color = "Random forest"))

plot_final <- plot_final + ggtitle("Error rate vs number of trees")
plot_final

```



The error rate for the AdaBoost model are clearly going down when the number of trees increases. Finally the model arrives at an error rate below 7% when 100 trees are included in the model. For the randomforest the pattern is less obvious, the error rate seems to go up and down as the number of trees in the model increases. 50 trees result in the lowest error rate. This error rate is also lower than the error rate produced by the best Adaboost model (100 trees). Therefore, for this spam classification, a randomforest with 50 trees

seems to be most suitable.

Question 2 Mixture Models

```
my_own_em <- function(K){  
  # 2 - Mixture Models ####  
  set.seed(1234567890)  
  
  max_it <- 100 # max number of EM iterations  
  min_change <- 0.1 # min change in log likelihood between two consecutive EM iterations  
  N=1000 # number of training points  
  D=10 # number of dimensions  
  x <- matrix(nrow=N, ncol=D) # training data  
  
  true_pi <- vector(length = K) # true mixing coefficients  
  true_mu <- matrix(nrow=K, ncol=D) # true conditional distributions  
  true_pi=c(rep(1/3, K))  
  
  if (K == 2){  
    true_mu[1,]=c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)  
    true_mu[2,]=c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)  
    plot(true_mu[1,], type="o", col="blue", ylim=c(0,1))  
    points(true_mu[2,], type="o", col="red")  
  }else if (K == 3){  
    true_mu[1,]=c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)  
    true_mu[2,]=c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)  
    true_mu[3,]=c(0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5)  
    plot(true_mu[1,], type="o", col="blue", ylim=c(0,1))  
    points(true_mu[2,], type="o", col="red")  
    points(true_mu[3,], type="o", col="green")  
  }else{  
    true_mu[1,]=c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)  
    true_mu[2,]=c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)  
    true_mu[3,]=c(0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5)  
    true_mu[4,]=c(0.3,0.5,0.5,0.7,0.5,0.5,0.5,0.5,0.4,0.5)  
    plot(true_mu[1,], type="o", col="blue", ylim=c(0,1))  
    points(true_mu[2,], type="o", col="red")  
    points(true_mu[3,], type="o", col="green")  
    points(true_mu[4,], type="o", col="yellow")  
  }  
  
  # Producing the training data  
  for(n in 1:N) {  
    k <- sample(1:K,1,prob=true_pi)  
    for(d in 1:D) {  
      x[n,d] <- rbinom(1,1,true_mu[k,d])  
    }  
  }  
  
  # number of guessed components
```

```

z <- matrix(nrow=N, ncol=K) # fractional component assignments
pi <- vector(length = K) # mixing coefficients
mu <- matrix(nrow=K, ncol=D) # conditional distributions
llik <- vector(length = max_it) # log likelihood of the EM iterations
# Random initialization of the parameters
pi <- runif(K,0.49,0.51)
pi <- pi / sum(pi)
for(k in 1:K) {
  mu[k,] <- runif(D,0.49,0.51)
}
pi
mu
for(it in 1:max_it) {
  if (K == 2){
    plot(mu[1,], type="o", col="blue", ylim=c(0,1))
    points(mu[2,], type="o", col="red")
  }else if (K == 3){
    plot(mu[1,], type="o", col="blue", ylim=c(0,1))
    points(mu[2,], type="o", col="red")
    points(mu[3,], type="o", col="green")
  }else{
    plot(mu[1,], type="o", col="blue", ylim=c(0,1))
    points(mu[2,], type="o", col="red")
    points(mu[3,], type="o", col="green")
    points(mu[4,], type="o", col="yellow")
  }
  Sys.sleep(0.5)
  # E-step: Computation of the fractional component assignments
  m <- matrix(NA, nrow = 1000, ncol = k)

  #Here I create the Bernoulli probabilities, lecture 1b, slide 7. I use 3 loops to do it for the three
  # not very efficient, but it works.
  for (j in 1:k){
    for(each in 1:nrow(x)){
      row <- x[each,]
      vec <- c()
      for (i in 1:10) {
        a <- mu[j,i]^row[i]
        b <- a * ((1-mu[j,i])^(1-row[i]))
        vec[i] <- b
        c <- prod(vec)
      }
      m[each, j] <- c
    }
  }

  # Here I create a empty matrix, to store all values for the numerator of the formula on the bottom of
  # slide 9, lecture 1b.
  m2 <- matrix(NA, ncol = k, nrow = 1000)

  # m2 stores all the values for the numerator of the formula on the bottom of slide 9, lecture 1b.
  for (i in 1:1000){
    a <- pi * m[i,]
  }

```

```

    m2[i,] <- a
  }

  # Sum m2 to get the denominator of the formula on the bottom of slide 9, lecture 1b.
  m2_sum <- rowSums(m2)
  m_final <- m2 / m2_sum

  #Log likelihood computation.
  ll <- matrix(nrow = 1000, ncol = K)
  for (j in 1:K){
    for (i in 1:1000){
      ll[i, j] <- sum(((x[i,] * log(mu[j,])) + (1 - x[i,])*log(1-mu[j,])))
    }
  }

  ll <- ll + pi
  llnew <- m_final * ll
  llik[it] <- sum(rowSums(llnew))

  cat("iteration: ", it, "log likelihood: ", llik[it], "\n")
  flush.console()
  # Stop if the log likelihood has not changed significantly
  if (it != 1){
    if (abs(llik[it] - llik[it-1]) < min_change) {break}
  }
  #M-step: ML parameter estimation from the data and fractional component assignments

  # Create the numerator for pi, slide 9, lecture 1b.
  numerator_pi <- colSums(m_final)

  # Create new values for pi, stored in the vector pi_new
  pi_new <- numerator_pi / N
  pi_new
  mnew <- matrix(NA, nrow = 1000, ncol = 10)
  mu_new <- matrix(NA, nrow = K, ncol = 10)

  for (j in 1:k){
    for (i in 1:1000){
      row <- x[i,] * m_final[i,j]
      mnew[i,] <- row
    }
    mnewsum <- colSums(mnew)/numerator_pi[j]
    mu_new[j,] <- mnewsum
  }

  # Now, to create the iterations, I have to run the code again and again, and specifying mu as new the
  # created for mu. Same goes for the other variables.
  mu <- mu_new
  pi <- pi_new
}
z <- m_final
output1 <- pi
output2 <- mu



```

```


output3 <- plot(lik[1:it], type="o")
z
result <- list(c(output1, output2, output3))
return(result)
}

```


```
my_own_em(2)
```




```
## iteration: 1 log likelihood: -6430.751
```




iteration: 2 log likelihood: -6417.599




iteration: 3 log likelihood: -6270.298




iteration: 4 log likelihood: -5381.969




iteration: 5 log likelihood: -4538.463




iteration: 6 log likelihood: -4463.134




iteration: 7 log likelihood: -4455.903




iteration: 8 log likelihood: -4453.165




iteration: 9 log likelihood: -4451.653




iteration: 10 log likelihood: -4450.674




iteration: 11 log likelihood: -4449.981




iteration: 12 log likelihood: -4449.461




iteration: 13 log likelihood: -4449.057




iteration: 14 log likelihood: -4448.734




iteration: 15 log likelihood: -4448.47




iteration: 16 log likelihood: -4448.251




iteration: 17 log likelihood: -4448.068




iteration: 18 log likelihood: -4447.913




iteration: 19 log likelihood: -4447.781



iteration: 20 log likelihood: -4447.669





iteration: 21 log likelihood: -4447.571




```
## [[1]]
## [1] 0.511053119 0.488946881 0.493173475 0.498954318 0.397460637
## [6] 0.625582274 0.596781124 0.380436306 0.278547959 0.717147834
## [11] 0.692791708 0.323034348 0.218495730 0.777869929 0.801849083
## [16] 0.204955853 0.111647688 0.914091323 0.880544386 0.089979192
## [21] 0.004290353 0.999714736
```

```
my_own_em(3)
```







iteration: 1 log likelihood: -6597.778




iteration: 2 log likelihood: -6595.239




iteration: 3 log likelihood: -6592.753




iteration: 4 log likelihood: -6573.7




iteration: 5 log likelihood: -6446.022




iteration: 6 log likelihood: -5978.865




iteration: 7 log likelihood: -5537.074




iteration: 8 log likelihood: -5429.225




iteration: 9 log likelihood: -5401.95




iteration: 10 log likelihood: -5389.023




iteration: 11 log likelihood: -5380.443




iteration: 12 log likelihood: -5373.845




iteration: 13 log likelihood: -5368.41




iteration: 14 log likelihood: -5363.759




iteration: 15 log likelihood: -5359.682




iteration: 16 log likelihood: -5356.051




iteration: 17 log likelihood: -5352.782




iteration: 18 log likelihood: -5349.816




iteration: 19 log likelihood: -5347.113




iteration: 20 log likelihood: -5344.641




iteration: 21 log likelihood: -5342.375




iteration: 22 log likelihood: -5340.295




iteration: 23 log likelihood: -5338.385




iteration: 24 log likelihood: -5336.63




iteration: 25 log likelihood: -5335.015




iteration: 26 log likelihood: -5333.529




iteration: 27 log likelihood: -5332.16




iteration: 28 log likelihood: -5330.9




iteration: 29 log likelihood: -5329.738




iteration: 30 log likelihood: -5328.666




iteration: 31 log likelihood: -5327.676




iteration: 32 log likelihood: -5326.762




iteration: 33 log likelihood: -5325.917




iteration: 34 log likelihood: -5325.135




iteration: 35 log likelihood: -5324.41




iteration: 36 log likelihood: -5323.739




iteration: 37 log likelihood: -5323.115




iteration: 38 log likelihood: -5322.537




iteration: 39 log likelihood: -5321.999




iteration: 40 log likelihood: -5321.498




iteration: 41 log likelihood: -5321.031




iteration: 42 log likelihood: -5320.596




iteration: 43 log likelihood: -5320.19




iteration: 44 log likelihood: -5319.81




iteration: 45 log likelihood: -5319.454




iteration: 46 log likelihood: -5319.121




iteration: 47 log likelihood: -5318.809




iteration: 48 log likelihood: -5318.515




iteration: 49 log likelihood: -5318.239




iteration: 50 log likelihood: -5317.979




iteration: 51 log likelihood: -5317.734




iteration: 52 log likelihood: -5317.503




iteration: 53 log likelihood: -5317.284




iteration: 54 log likelihood: -5317.077




iteration: 55 log likelihood: -5316.881




iteration: 56 log likelihood: -5316.695




iteration: 57 log likelihood: -5316.518




iteration: 58 log likelihood: -5316.349




iteration: 59 log likelihood: -5316.189




iteration: 60 log likelihood: -5316.036




iteration: 61 log likelihood: -5315.89




iteration: 62 log likelihood: -5315.75




iteration: 63 log likelihood: -5315.616




iteration: 64 log likelihood: -5315.487




iteration: 65 log likelihood: -5315.364




iteration: 66 log likelihood: -5315.246




iteration: 67 log likelihood: -5315.132




iteration: 68 log likelihood: -5315.022




iteration: 69 log likelihood: -5314.916



iteration: 70 log likelihood: -5314.814





```
## iteration: 71 log likelihood: -5314.715
```




```
## [[1]]
## [1] 0.32411556 0.30717327 0.36871116 0.47383554 0.49067297 0.50907301
## [7] 0.38106095 0.47962547 0.58355730 0.62860603 0.47050854 0.41965465
## [13] 0.30784064 0.47936473 0.71594258 0.69512193 0.53269167 0.29039403
## [19] 0.19764738 0.49158731 0.76736739 0.78836950 0.46597001 0.23146477
## [25] 0.13424145 0.48919757 0.85221677 0.89207742 0.49348782 0.10652662
## [31] 0.01757154 0.39918885 0.99992807
```


```
my_own_em(4)
```




iteration: 1 log likelihood: -6680.657




iteration: 2 log likelihood: -6654.874




iteration: 3 log likelihood: -6650.741




iteration: 4 log likelihood: -6628.679




iteration: 5 log likelihood: -6525.998




iteration: 6 log likelihood: -6241.218




iteration: 7 log likelihood: -5962.695




iteration: 8 log likelihood: -5852.656




iteration: 9 log likelihood: -5804.099




iteration: 10 log likelihood: -5774.053




iteration: 11 log likelihood: -5754.55




iteration: 12 log likelihood: -5741.968




iteration: 13 log likelihood: -5733.572




iteration: 14 log likelihood: -5727.573




iteration: 15 log likelihood: -5722.953




iteration: 16 log likelihood: -5719.164




iteration: 17 log likelihood: -5715.91




iteration: 18 log likelihood: -5713.019




iteration: 19 log likelihood: -5710.383




iteration: 20 log likelihood: -5707.926




iteration: 21 log likelihood: -5705.593




iteration: 22 log likelihood: -5703.335




iteration: 23 log likelihood: -5701.111




iteration: 24 log likelihood: -5698.879




iteration: 25 log likelihood: -5696.603




iteration: 26 log likelihood: -5694.243




iteration: 27 log likelihood: -5691.767




iteration: 28 log likelihood: -5689.141




iteration: 29 log likelihood: -5686.341




iteration: 30 log likelihood: -5683.349




iteration: 31 log likelihood: -5680.162




iteration: 32 log likelihood: -5676.789




iteration: 33 log likelihood: -5673.256




iteration: 34 log likelihood: -5669.605




iteration: 35 log likelihood: -5665.886




iteration: 36 log likelihood: -5662.158




iteration: 37 log likelihood: -5658.48




iteration: 38 log likelihood: -5654.903




iteration: 39 log likelihood: -5651.468




iteration: 40 log likelihood: -5648.201




iteration: 41 log likelihood: -5645.114




iteration: 42 log likelihood: -5642.206




iteration: 43 log likelihood: -5639.468




iteration: 44 log likelihood: -5636.879




iteration: 45 log likelihood: -5634.418




iteration: 46 log likelihood: -5632.056




iteration: 47 log likelihood: -5629.768




iteration: 48 log likelihood: -5627.526




iteration: 49 log likelihood: -5625.304




iteration: 50 log likelihood: -5623.08




iteration: 51 log likelihood: -5620.832




iteration: 52 log likelihood: -5618.543




iteration: 53 log likelihood: -5616.204




iteration: 54 log likelihood: -5613.809




iteration: 55 log likelihood: -5611.36




iteration: 56 log likelihood: -5608.87




iteration: 57 log likelihood: -5606.356




iteration: 58 log likelihood: -5603.846




iteration: 59 log likelihood: -5601.369




iteration: 60 log likelihood: -5598.96




iteration: 61 log likelihood: -5596.652




iteration: 62 log likelihood: -5594.475




iteration: 63 log likelihood: -5592.454




iteration: 64 log likelihood: -5590.605




iteration: 65 log likelihood: -5588.94




iteration: 66 log likelihood: -5587.461




iteration: 67 log likelihood: -5586.164




iteration: 68 log likelihood: -5585.04




iteration: 69 log likelihood: -5584.076




iteration: 70 log likelihood: -5583.259




iteration: 71 log likelihood: -5582.571




iteration: 72 log likelihood: -5581.998




iteration: 73 log likelihood: -5581.523




iteration: 74 log likelihood: -5581.132




iteration: 75 log likelihood: -5580.812




iteration: 76 log likelihood: -5580.552




iteration: 77 log likelihood: -5580.341




iteration: 78 log likelihood: -5580.169




iteration: 79 log likelihood: -5580.029



iteration: 80 log likelihood: -5579.915



iteration: 81 log likelihood: -5579.821



```
## [[1]]
## [1] 0.26041646 0.25364281 0.29318139 0.19275935 0.38658629 0.52163416
## [7] 0.44004028 0.32167671 0.39321260 0.61225557 0.40500112 0.71886505
## [13] 0.48795286 0.38916318 0.54876590 0.53607077 0.57710846 0.71360156
## [19] 0.33179850 0.61441593 0.50636776 0.27235697 0.65703184 0.51059837
## [25] 0.51395263 0.77891604 0.20540033 0.46364689 0.24003794 0.21316485
## [31] 0.78592225 0.72631841 0.34159808 0.93237925 0.17009121 0.63127091
## [37] 0.37660293 0.08737204 0.80624121 0.43780078 0.47606925 0.99886761
## [43] 0.04609842 0.49363597
```

Appendix

```
# Loading packages and importing files ####
library(mboost)
library(randomForest)
library(ggplot2)
sp <- read.csv2("spambase.csv", header = FALSE, sep = ",", stringsAsFactors = FALSE)
num_sp <- data.frame(data.matrix(sp))
num_sp$V58 <- factor(num_sp$V58)
# shuffling data and dividing into train and test ####
n <- dim(num_sp)[1]
ncol <- dim(num_sp)[2]
set.seed(1234567890)
id <- sample(1:n, floor(n*(2/3)))
train <- num_sp[id,]
test <- num_sp[-id,]
# Adaboost
ntree <- c(10, 20, 30, 40, 50, 60, 70, 80, 90, 100)
error <- c()
```



```

for (i in seq(from = 10, to = 100, by = 10)){
  bb <- blackboost(V58 ~., data = train, control = boost_control(mstop = i), family = AdaExp())
  bb_predict <- predict(bb, newdata = test, type = c("class"))
  confusion_bb <- table(test$V58, bb_predict)
  miss_class_bb <- (confusion_bb[1,2] + confusion_bb[2,1])/nrow(test)
  error[(i/10)] <- miss_class_bb
}

error_df <- data.frame(cbind(ntree, error))
# Random forest #####
ntree_rf <- c(10, 20, 30, 40, 50, 60, 70, 80, 90, 100)
error_rf <- c()

for (i in seq(from = 10, to = 100, by = 10)){
  rf <- randomForest(V58 ~., data = train, ntree= 10)
  rf_predict <- predict(rf, newdata = test, type = c("class"))
  confusion_rf <- table(test$V58, rf_predict)
  miss_class_rf <- (confusion_rf[1,2] + confusion_rf[2,1])/nrow(test)
  error_rf[i/10] <- miss_class_rf
}

error_df_rf <- data.frame(cbind(ntree_rf, error_rf))

df <- cbind(error_df, error_df_rf)
df <- df[, -3]

plot_final <- ggplot(df, aes(ntree)) +
  geom_line(aes(y=error, color = "Adaboost")) +
  geom_line(aes(y=error_rf, color = "Random forest"))

plot_final <- plot_final + ggtitle("Error rate vs number of trees")
plot_final

my_own_em <- function(K){
  # 2 - Mixture Models #####
  set.seed(1234567890)

  max_it <- 100 # max number of EM iterations
  min_change <- 0.1 # min change in log likelihood between two consecutive EM iterations
  N=1000 # number of training points
  D=10 # number of dimensions
  x <- matrix(nrow=N, ncol=D) # training data

  true_pi <- vector(length = K) # true mixing coefficients
  true_mu <- matrix(nrow=K, ncol=D) # true conditional distributions
  true_pi=c(rep(1/3, K))

  if (K == 2){
    true_mu[1,]=c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)
    true_mu[2,]=c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
    plot(true_mu[1,], type="o", col="blue", ylim=c(0,1))
    points(true_mu[2,], type="o", col="red")
  }
}

```

```

}else if (K == 3){
  true_mu[1,]=c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)
  true_mu[2,]=c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
  true_mu[3,]=c(0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5)
  plot(true_mu[1,], type="o", col="blue", ylim=c(0,1))
  points(true_mu[2,], type="o", col="red")
  points(true_mu[3,], type="o", col="green")
}else{
  true_mu[1,]=c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)
  true_mu[2,]=c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
  true_mu[3,]=c(0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5)
  true_mu[4,]=c(0.3,0.5,0.5,0.7,0.5,0.5,0.5,0.5,0.4,0.5)
  plot(true_mu[1,], type="o", col="blue", ylim=c(0,1))
  points(true_mu[2,], type="o", col="red")
  points(true_mu[3,], type="o", col="green")
  points(true_mu[4,], type="o", col="yellow")
}

# Producing the training data
for(n in 1:N) {
  k <- sample(1:K,1,prob=true_pi)
  for(d in 1:D) {
    x[n,d] <- rbinom(1,1,true_mu[k,d])
  }
}

# number of guessed components
z <- matrix(nrow=N, ncol=K) # fractional component assignments
pi <- vector(length = K) # mixing coefficients
mu <- matrix(nrow=K, ncol=D) # conditional distributions
llik <- vector(length = max_it) # log likelihood of the EM iterations
# Random initialization of the paramters
pi <- runif(K,0.49,0.51)
pi <- pi / sum(pi)
for(k in 1:K) {
  mu[k,] <- runif(D,0.49,0.51)
}
pi
mu
for(it in 1:max_it) {
  if (K == 2){
    plot(mu[1,], type="o", col="blue", ylim=c(0,1))
    points(mu[2,], type="o", col="red")
  }else if (K == 3){
    plot(mu[1,], type="o", col="blue", ylim=c(0,1))
    points(mu[2,], type="o", col="red")
    points(mu[3,], type="o", col="green")
  }else{
    plot(mu[1,], type="o", col="blue", ylim=c(0,1))
    points(mu[2,], type="o", col="red")
    points(mu[3,], type="o", col="green")
    points(mu[4,], type="o", col="yellow")
  }
}

```

```

}
Sys.sleep(0.5)
# E-step: Computation of the fractional component assignments
m <- matrix(NA, nrow = 1000, ncol = k)

# Here I create the Bernoulli probabilities, lecture 1b, slide 7. I use 3 loops to do it for the three
# not very efficient, but it works.
for (j in 1:k){
  for(each in 1:nrow(x)){
    row <- x[each,]
    vec <- c()
    for (i in 1:10) {
      a <- mu[j,i]^row[i]
      b <- a * ((1-mu[j,i])^(1-row[i]))
      vec[i] <- b
      c <- prod(vec)
    }
    m[each, j] <- c
  }
}

# Here I create a empty matrix, to store all values for the numerator of the formula on the bottom of
# slide 9, lecture 1b.
m2 <- matrix(NA, ncol = k, nrow = 1000)

# m2 stores all the values for the numerator of the formula on the bottom of slide 9, lecture 1b.
for (i in 1:1000){
  a <- pi * m[i,]
  m2[i,] <- a
}

# Sum m2 to get the denominator of the formula on the bottom of slide 9, lecture 1b.
m2_sum <- rowSums(m2)
m_final <- m2 / m2_sum

#Log likelihood computation.
ll <- matrix(nrow = 1000, ncol = K)
for (j in 1:K){
  for (i in 1:1000){
    ll[i, j] <- sum((x[i,] * log(mu[j,])) + (1 - x[i,])*log(1-mu[j,])))
  }
}

ll <- ll + pi
llnew <- m_final * ll
llik[it] <- sum(rowSums(llnew))

cat("iteration: ", it, "log likelihood: ", llik[it], "\n")
flush.console()
# Stop if the log likelihood has not changed significantly
if (it != 1){
  if (abs(llik[it] - llik[it-1]) < min_change) {break}
}

```

```

#M-step: ML parameter estimation from the data and fractional component assignments

# Create the numerator for pi, slide 9, lecture 1b.
numerator_pi <- colSums(m_final)

# Create new values for pi, stored in the vector pi_new
pi_new <- numerator_pi / N
pi_new
mnew <- matrix(NA, nrow = 1000, ncol = 10)
mu_new <- matrix(NA, nrow = K, ncol = 10)

for (j in 1:k){
  for (i in 1:1000){
    row <- x[i,] * m_final[i,j]
    mnew[i,] <- row
  }
  mnewsum <- colSums(mnew)/numerator_pi[j]
  mu_new[j,] <- mnewsum
}

# Now, to create the iterations, I have to run the code again and again, and specifying mu as new the
# created for mu. Same goes for the other variables.
mu <- mu_new
pi <- pi_new
}
z <- m_final
output1 <- pi
output2 <- mu
output3 <- plot(llik[1:it], type="o")
z
result <- list(c(output1, output2, output3))
return(result)
}
my_own_em(2)
my_own_em(3)
my_own_em(4)

```